Object Detection

Nov. 26-27

What is object detection

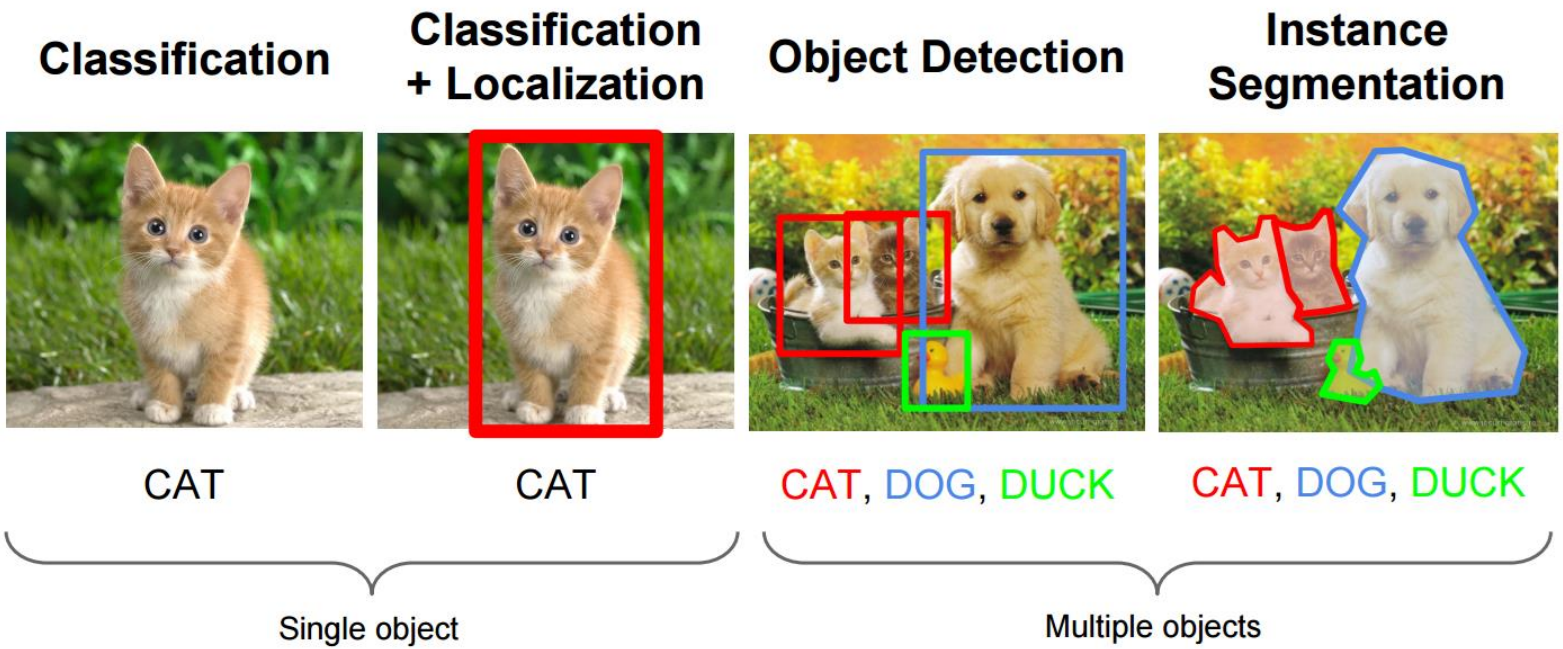

http://cs231n.stanford.edu/slides/winter1516_lecture8.pdf

Terms

Recall

Precision

mAP

IoU

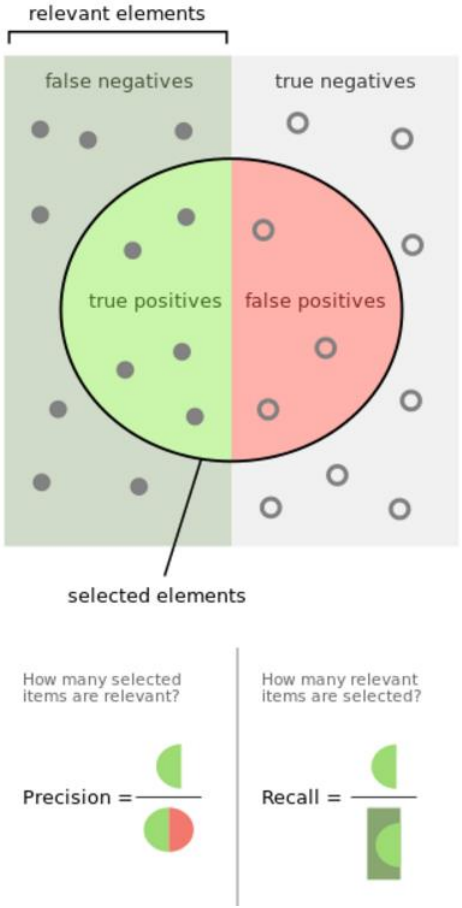

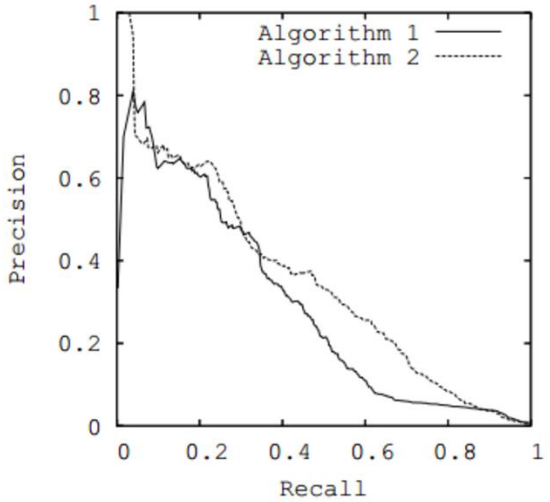

https://en.wikipedia.org/wiki/Precision_and_recall

Detection Competitions

Pascal VOC

COCO

ImageNet ILSVRC

VOC: 20 classes

COCO: 200 classes

http://host.robots.ox.ac.uk/pascal/VOC/voc2012/index.html#introduction

Before Deep Learning

Sliding windows.

- Score every subwindow.

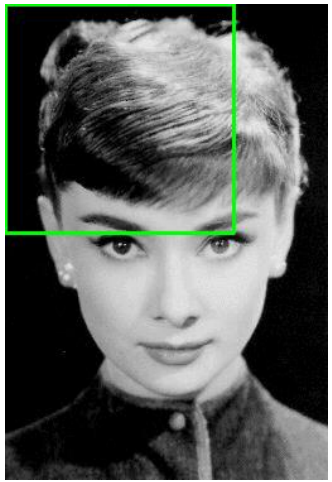

http://www.pyimagesearch.com/2014/11/10/histogram-oriented-gradients-object-detection/

Deformable part models (DPM)

- Uses HOG features
- Very fast

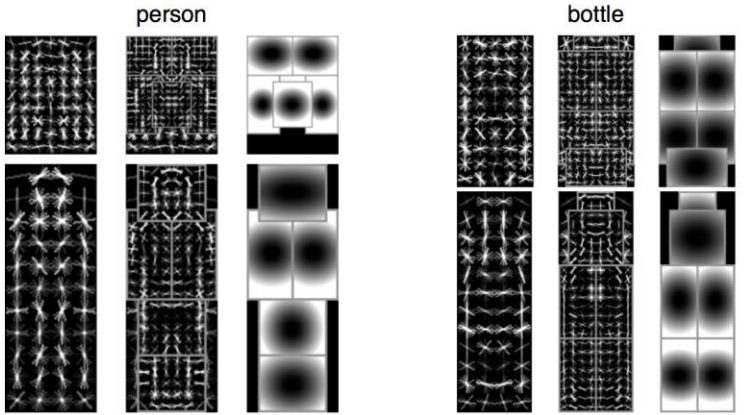

https://cs.brown.edu/~pff/papers/lsvm-pami.pdf

Selective Search

http://www.huppelen.nl/publications/selectiveSearchDraft.pdf

Hard Negative Mining

Imbalance between positive and negative examples.

Use negative examples with higher confidence score.

Non Maximum Suppression

If output boxes overlap, only consider the most confident.

Bounding Box Regression

Regression used to find bounding box parameters

Applied in one of two ways

- Bounding box refinement
- Complete object detection

Example Loss Function

$$\sum_{i \ in \ I} [(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 + (\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2]$$

I is the set of all matching bounding boxes (Highest IoU with ground truth)

Two Stage Detection

Part I

RCNN : Region Proposal + CNN

Use selective search to come up with regional proposal

First object detection method using CNN

Rich feature hierarchies for accurate object detection and semantic segmentation
Ross Girshick Jeff Donahue Trevor Darrell Jitendra Malik
Nov 2013

https://people.eecs.berkeley.edu/~rbg/papers/r-cnn-cvpr.pdf

RCNN :

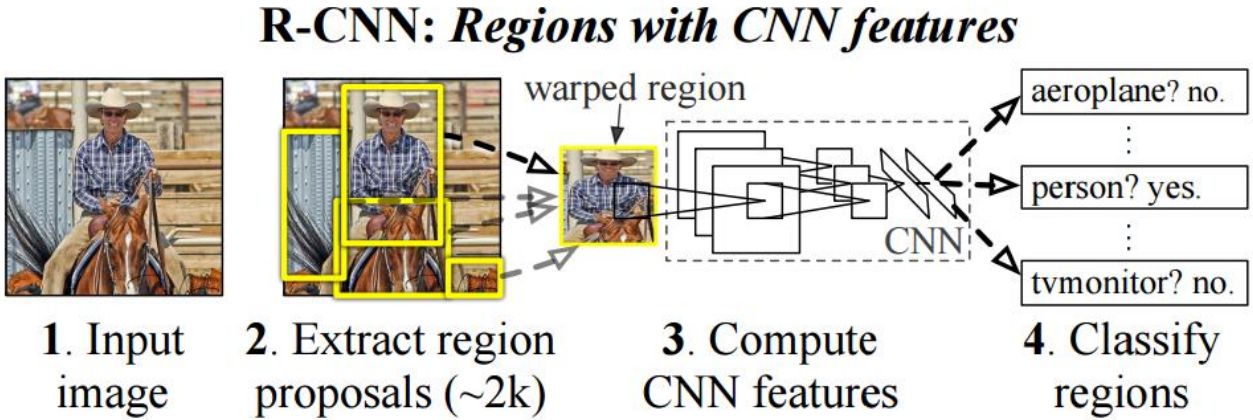

https://people.eecs.berkeley.edu/~rbg/papers/r-cnn-cvpr.pdf

Training RCNN

Step1: train your own CNN model for classification (or use existing model), using ImageNet dataset.

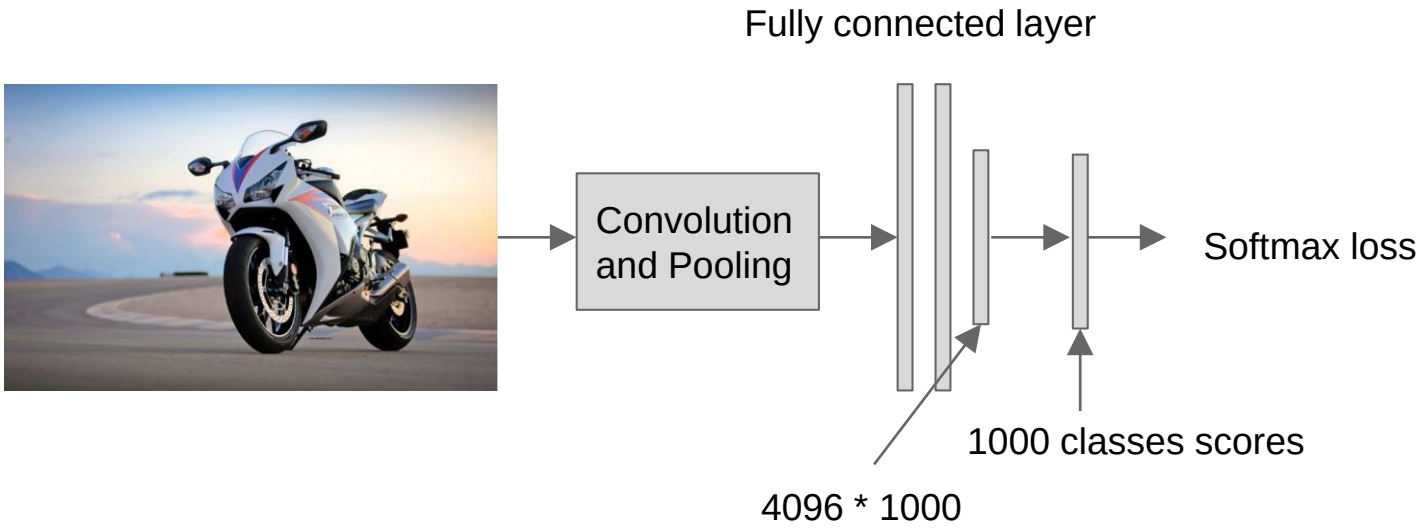

Image source : http://www.shunvmall.com/bike-pic/47539009.html

Training RCNN

Step2: focus on 20 classes + 1 background. Remove the last FC layer and replace it with a smaller layer and fine-tune the model using PASCAL VOC dataset

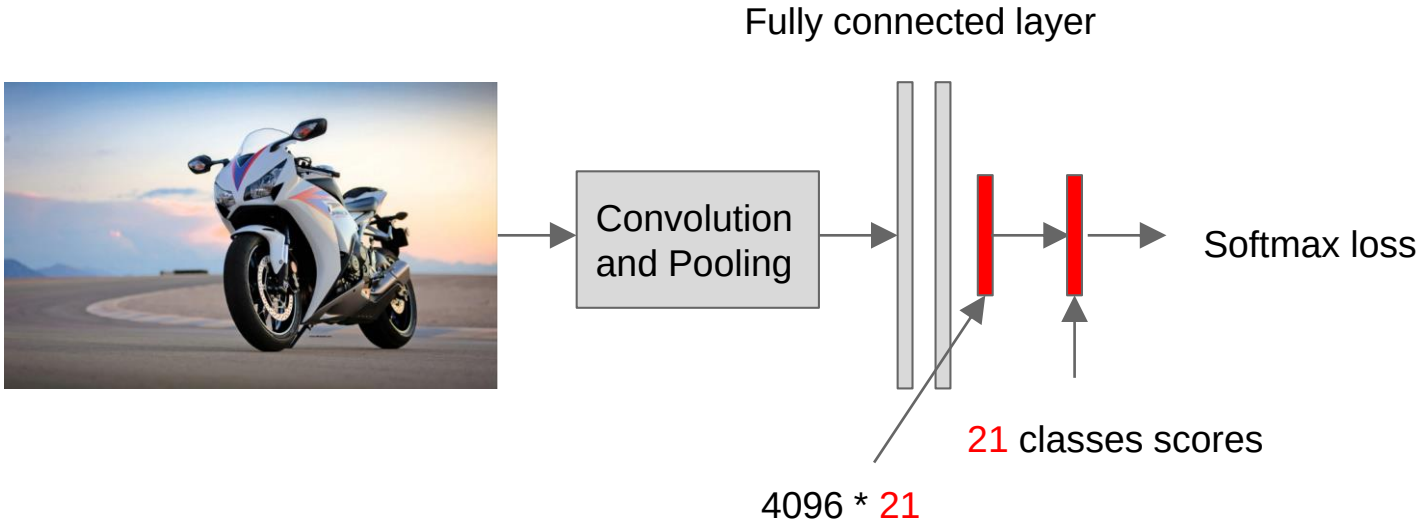

Training RCNN

Step3: extract feature

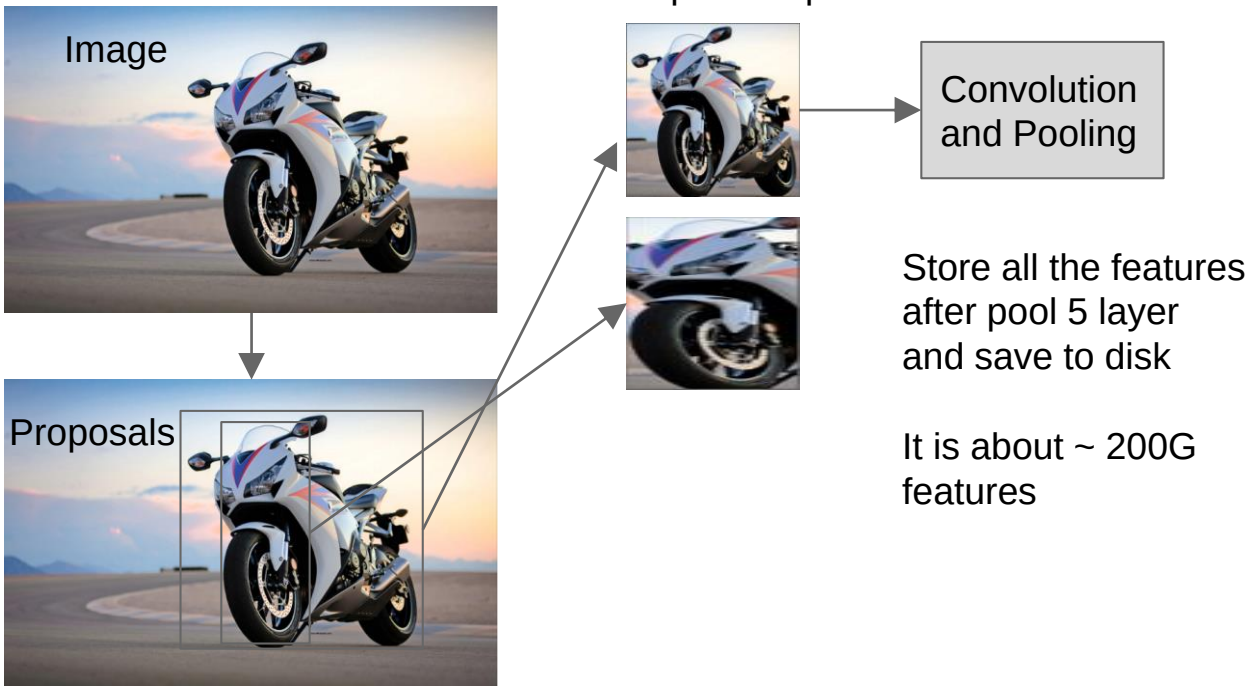

Training RCNN

Step4: train SVM for each class

Crop / Warp image

Features from last step

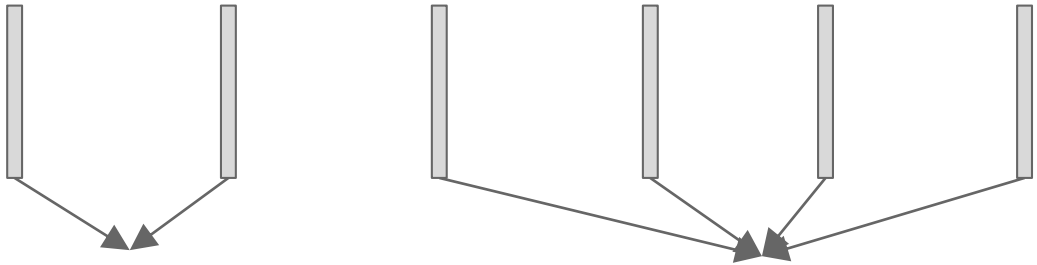

Positive samples for Motorbike

Negative samples for Motorbike

Training RCNN

Step4: train SVM for each class

Crop / Warp image

Features from last step

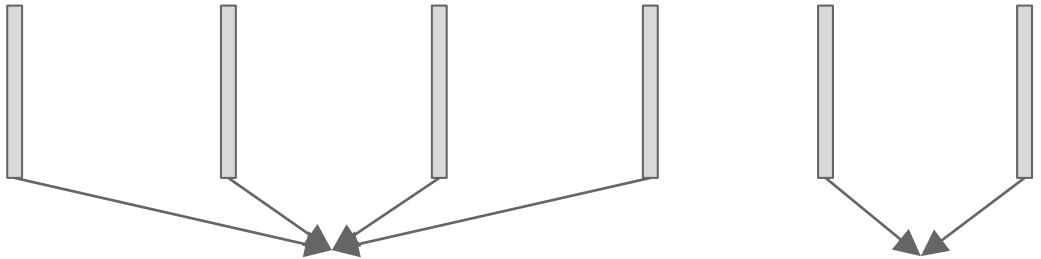

Negative samples for Bicycle

Positive samples for Bicycle

Fast RCNN

Share convolution layers for proposals from the same image

Faster and More accurate than RCNN

ROI Pooling

Fast R-CNN
Ross Girshick
Apr 2015

https://arxiv.org/pdf/1504.08083v2.pdf

Fast RCNN

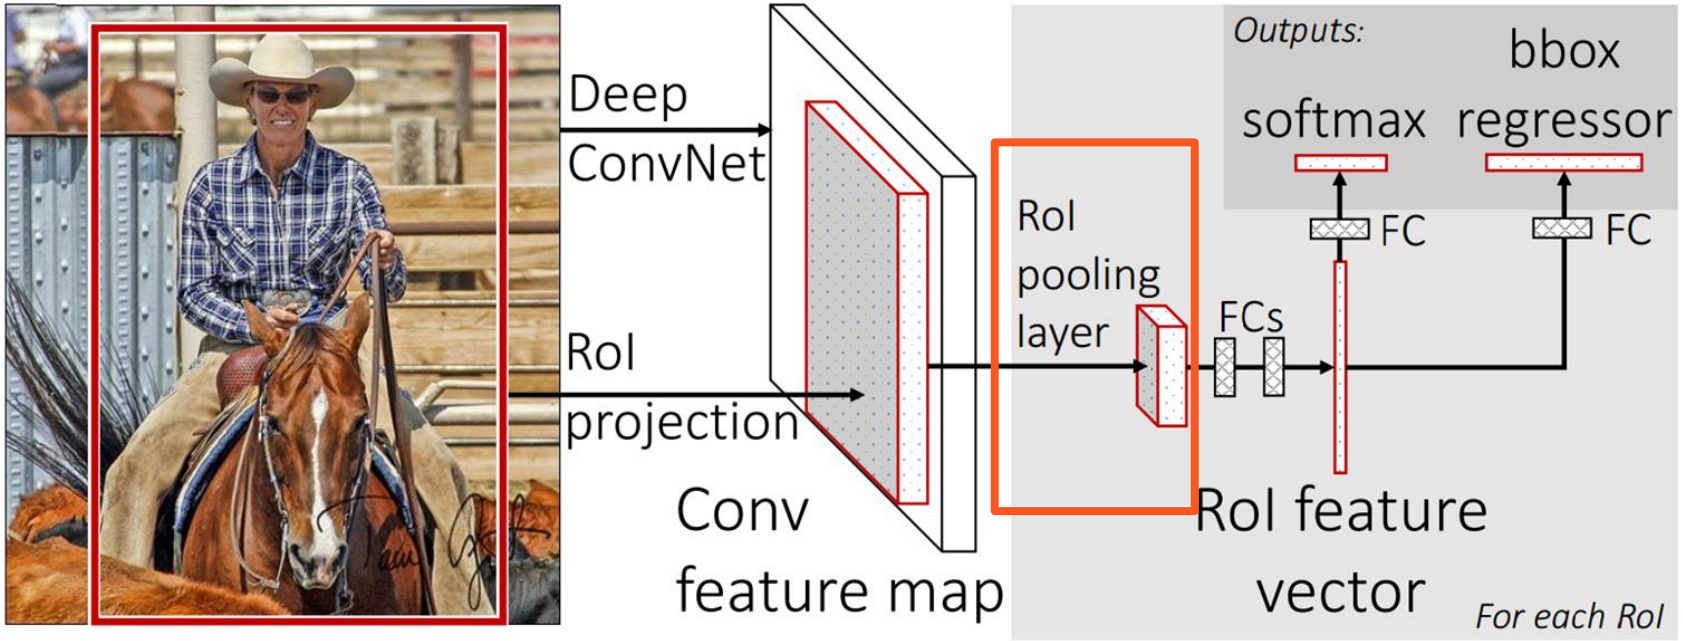

https://arxiv.org/pdf/1504.08083v2.pdf

RoI Pooling

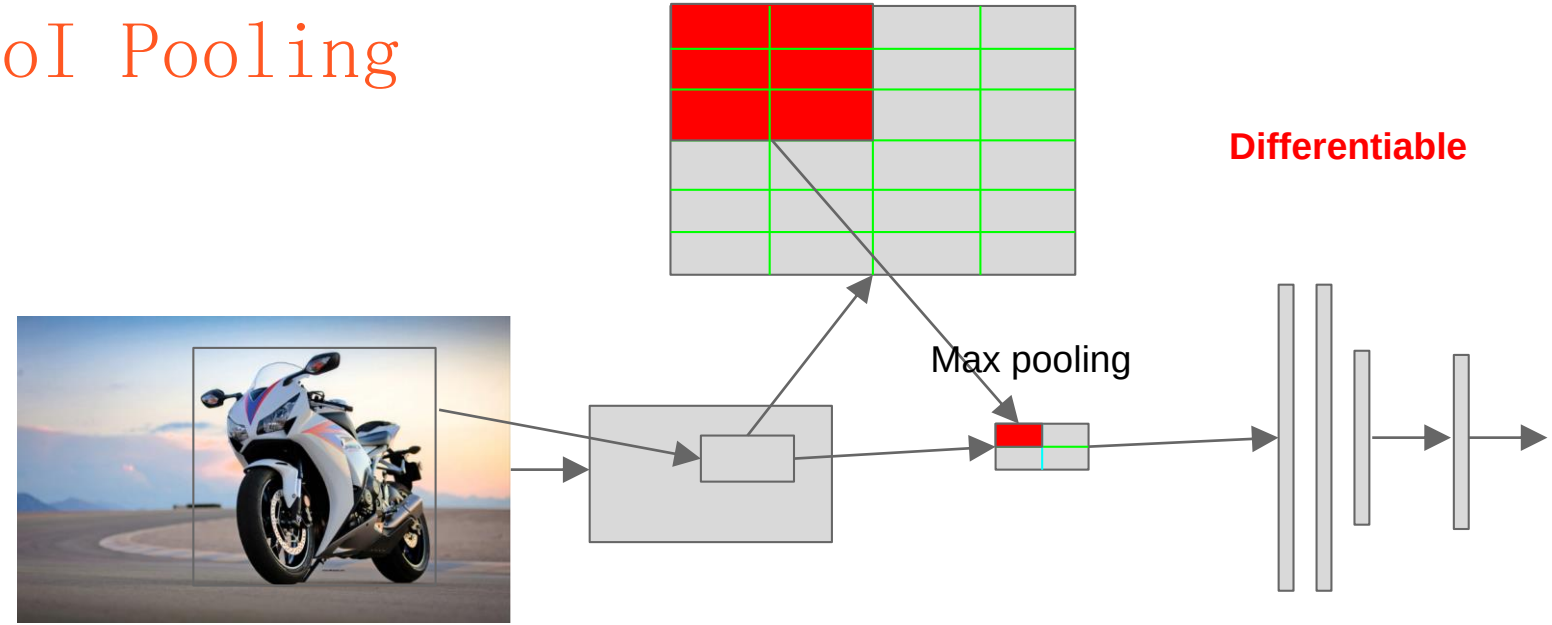
Differentiable
Max pooling
image
Conv layer
Divided into
h * w region

ROI Pooling

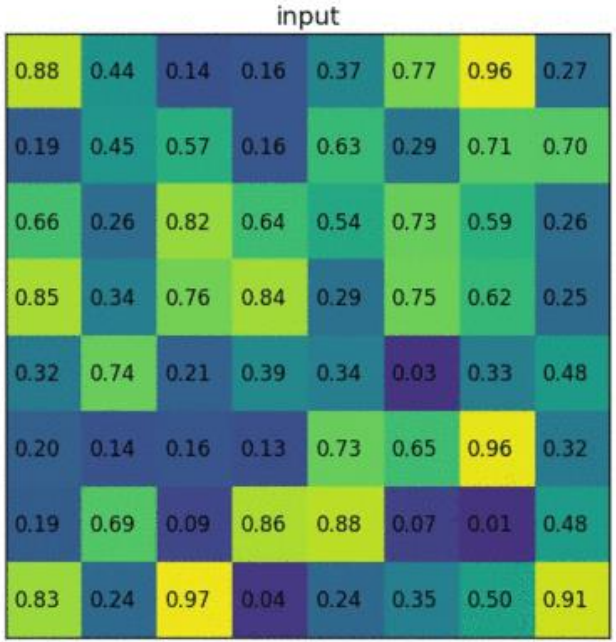

https://blog.deepsense.ai/region-of-interest-pooling-explained/

What is bbox-regressor?

Bounding box regression

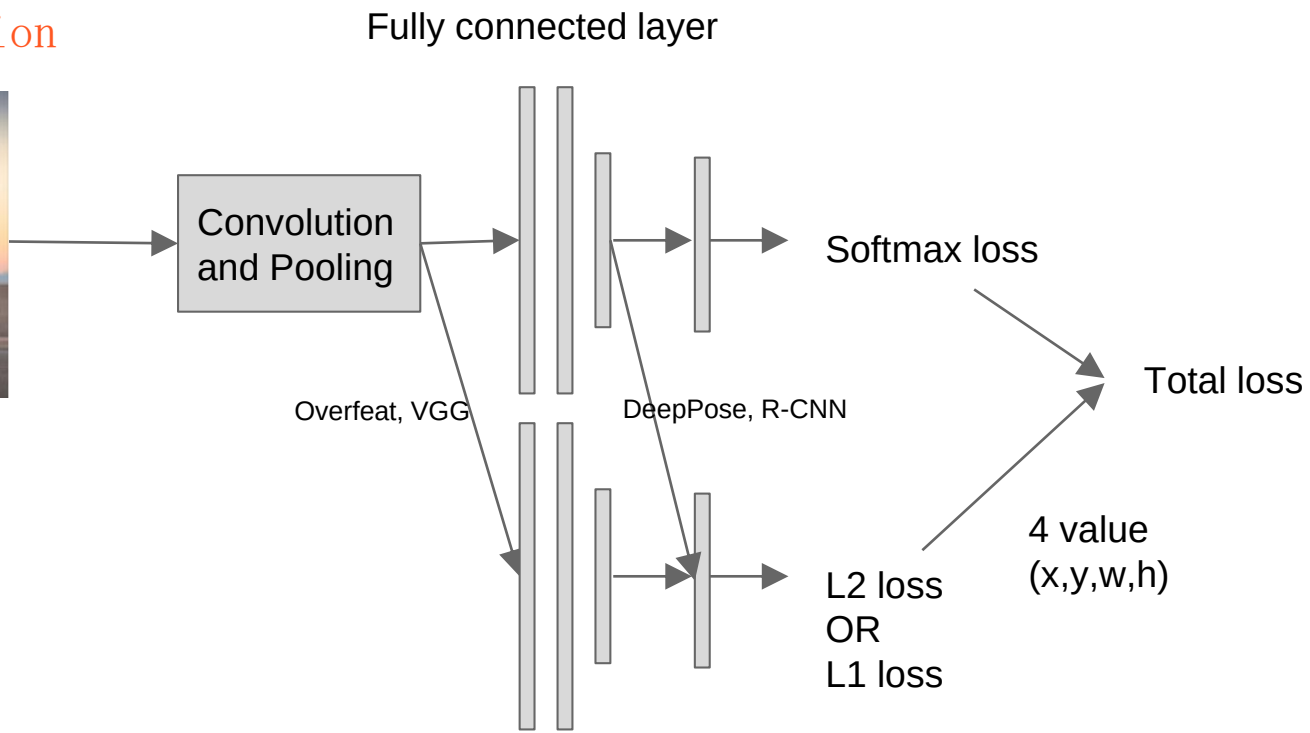

Convolution and Pooling

Fully connected layer

Softmax loss

Total loss

4 value (x,y,w,h)

Smooth L1 loss

Result compare

	Fast R-CNN	R-CNN
Train time (h)	9.5	84
-speedup	8.8x	1x
Test time/image	0.32s	47.00 s
-test speedup	146x	1x
mAP	66.9	66

Trained using VGG 16 on Pascal VOC 2007 dataset
Not including proposal time

Source: R. Girshick

Result compare

	Fast R-CNN	R-CNN
Test time/image	0.32s	47.00 s
-test speedup	146x	1x
Test time/image with proposal	2s	50 s
-test speedup	25x	1x

Source: cs231 standford

Faster RCNN

Don’t need to have external regional proposals

RPN - Regional Proposal Network

Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks

Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun

Jun 2015

https://arxiv.org/pdf/1506.01497v3.pdf

Faster RCNN

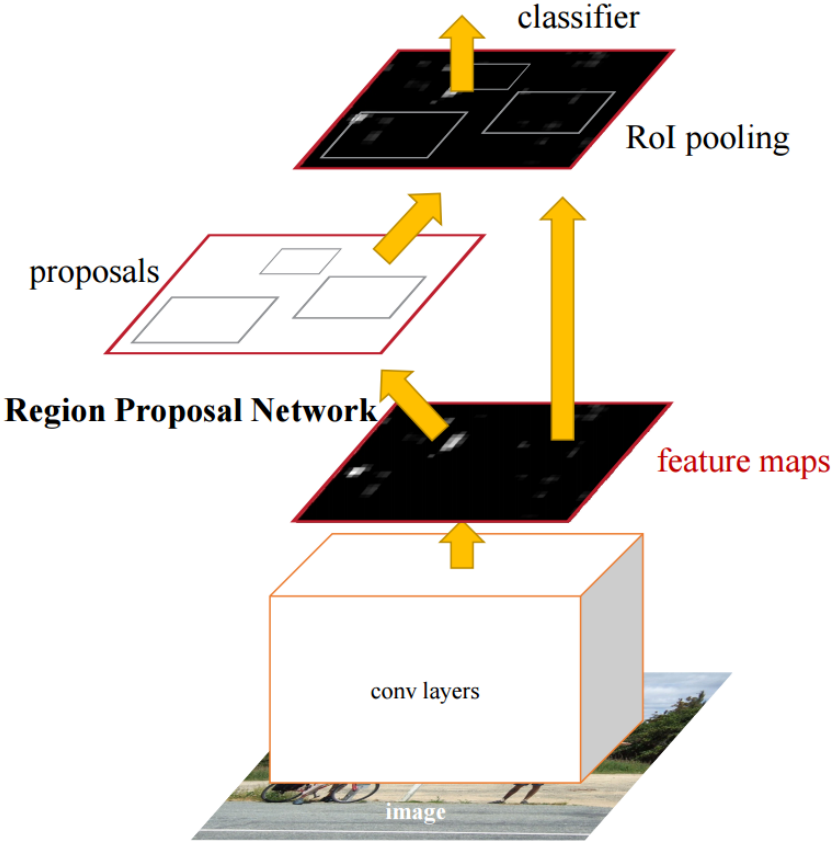

https://arxiv.org/pdf/1506.01497v3.pdf

Faster RCNN

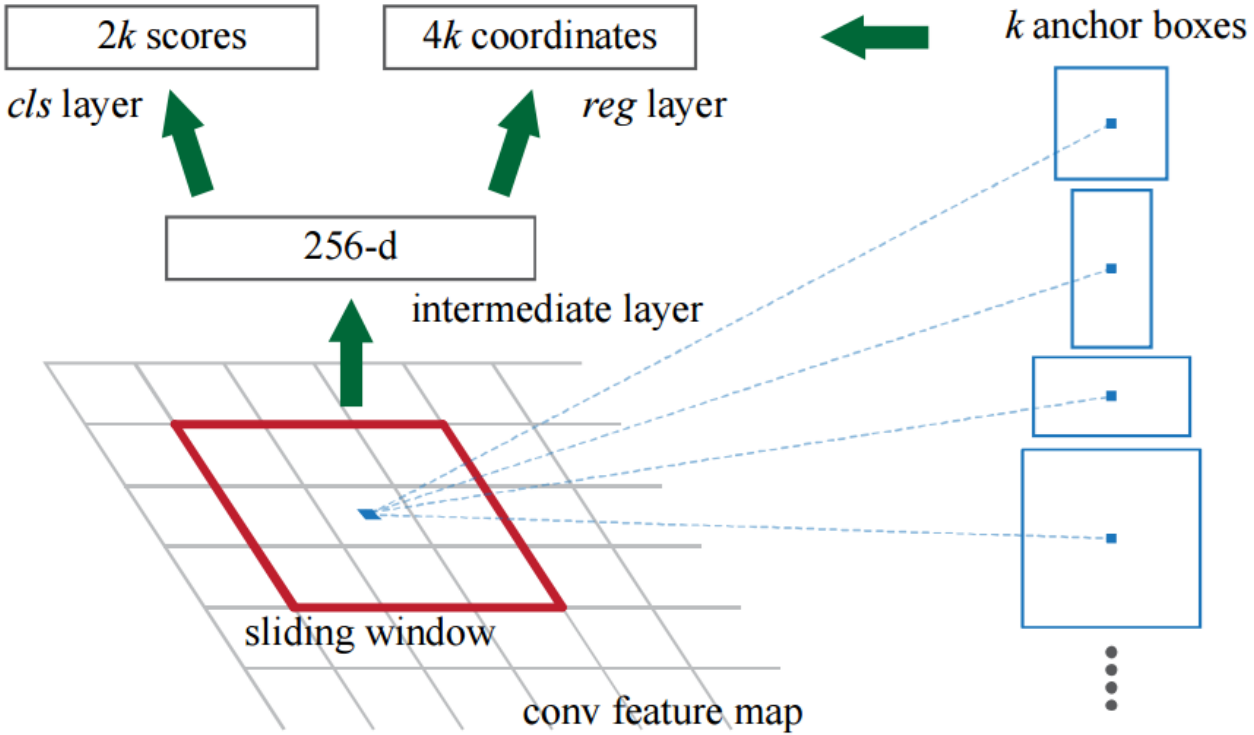

https://arxiv.org/pdf/1506.01497v3.pdf

Result compare

	Faster R-CNN	Fast R-CNN	R-CNN
Test time/image With proposal	0.2S	2s	50s
-test speedup	250x	25x	1x
mAP	66.9	66.9	66

Trained using Pascal VOC 2007 dataset

Source: cs231 standford

R-FCN :Region-based Fully Convolutional Networks

Use position sensitive score map

Share all conv and fc layers between all proposals for the same image

R-FCN: Object Detection via Region-based Fully Convolutional Networks
Jifeng Dai, Yi Li, Kaiming He, Jian Sun
May 2016

https://arxiv.org/pdf/1605.06409v2.pdf

R-FCN

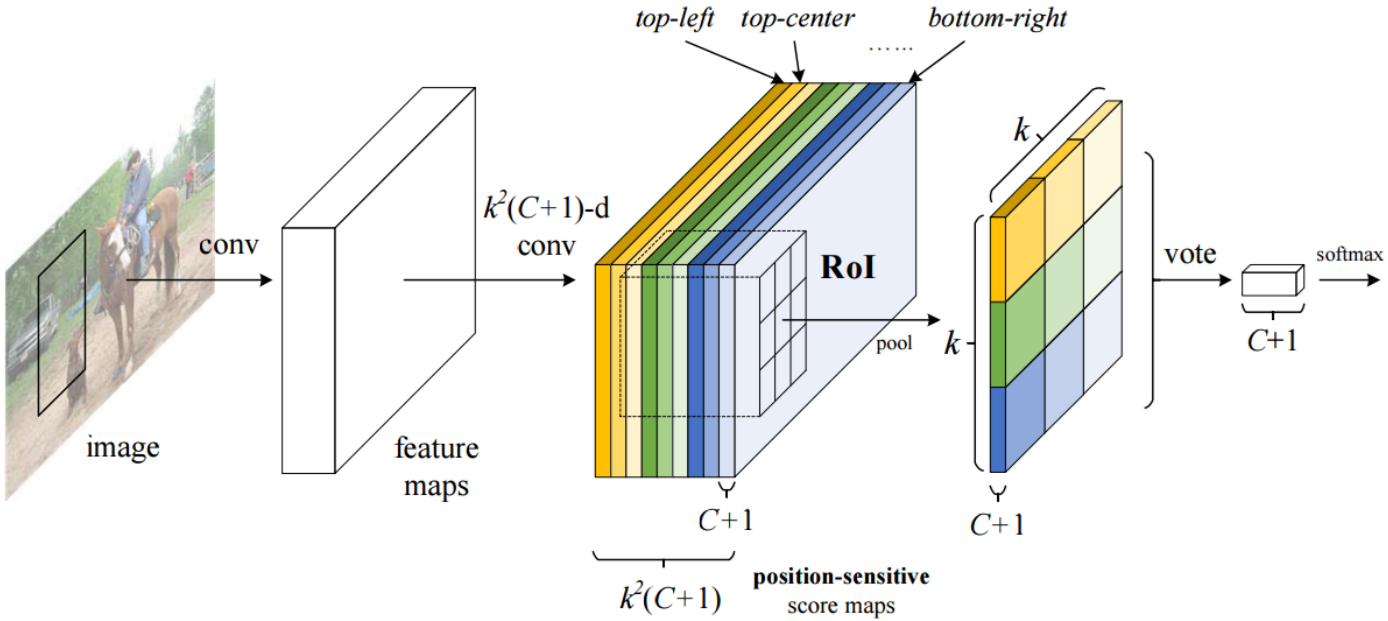

Figure 1: Key idea of **R-FCN** for object detection. In this illustration, there are $k \times k = 3 \times 3$ position-sensitive score maps generated by a fully convolutional network. For each of the $k \times k$ bins in an RoI, pooling is only performed on one of the k^2 maps (marked by different colors).

R-FCN

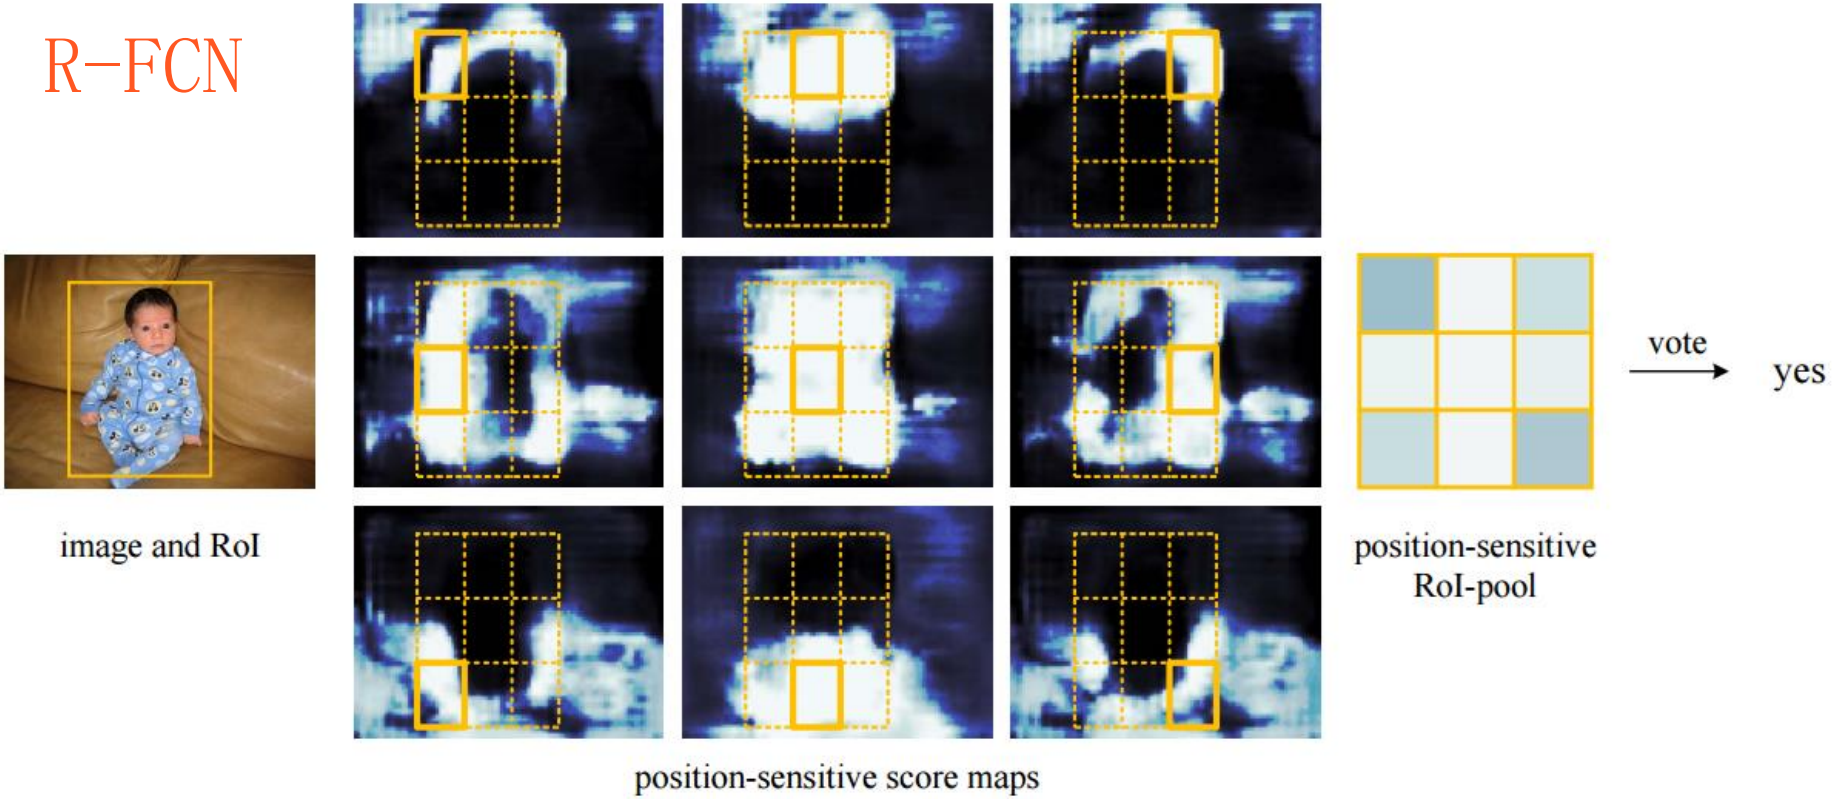

Figure 3: Visualization of R-FCN ($k \times k = 3 \times 3$) for the *person* category.

R-FCN

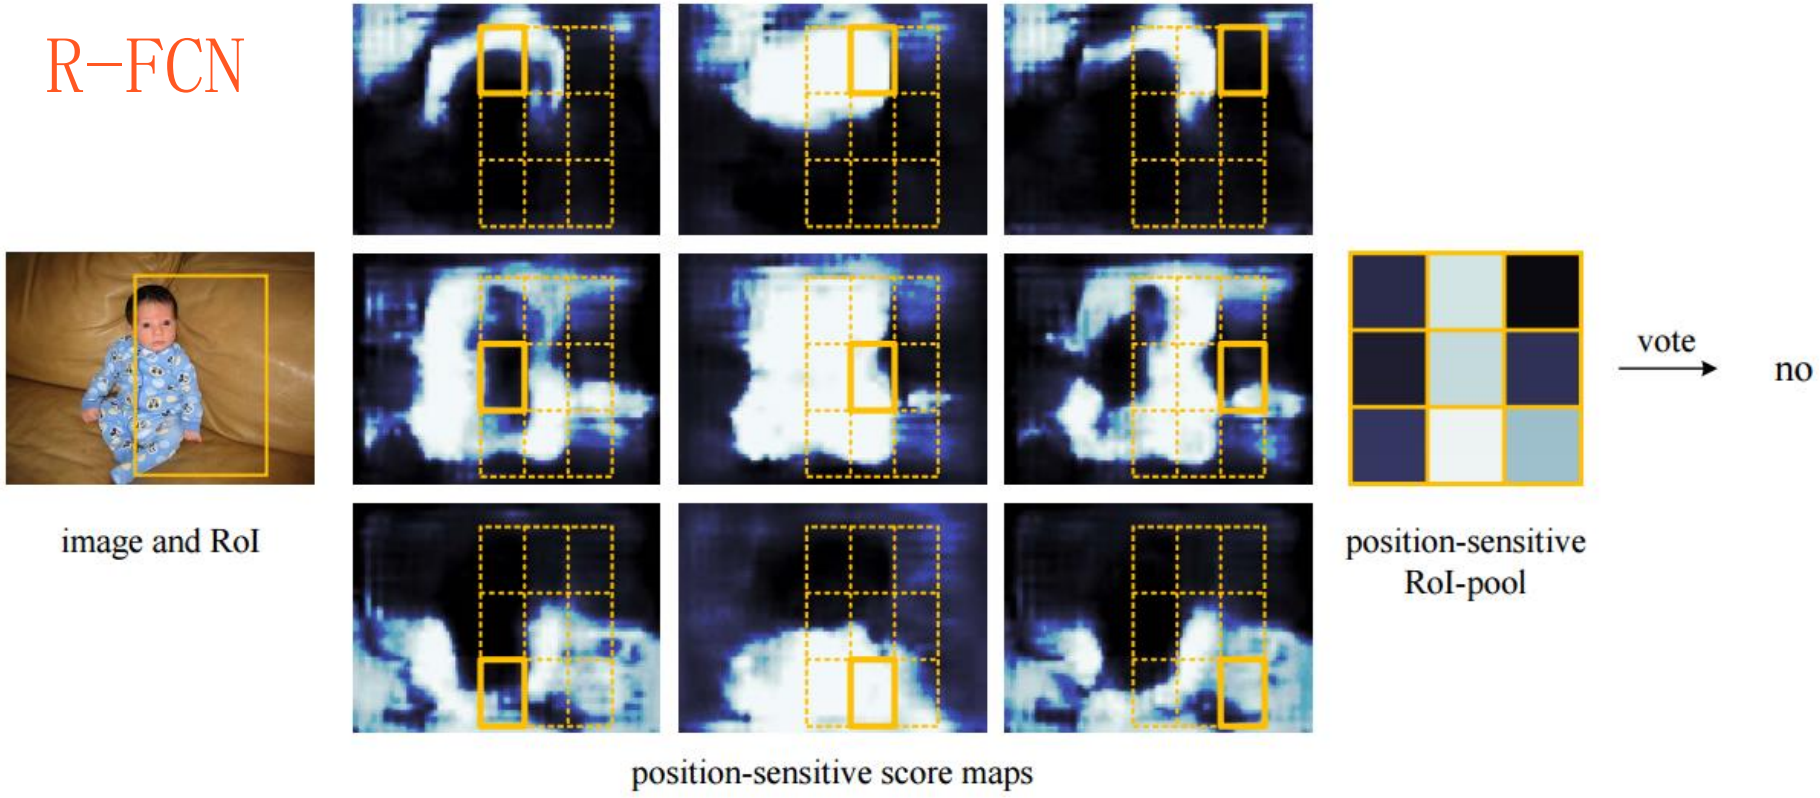

Figure 4: Visualization when an RoI does not correctly overlap the object.

Methodologies Compare

	RCNN	Faster RCNN	RFCN
Depth of shared convolutional subnetwork	0	91	101
Depth of ROI-wise subnetwork	101	10	0

Trained using ResNet 101

Result compare

	Faster R-CNN	R-FCN
Test time/image With proposal	0.42S	0.2s
mAP	76.4	76.6

Trained using ResNet 101 on Pascal VOC 2007 dataset

Unified Detection

Part II

Problems with 2 step detection.

Complex Pipeline

Slow (Cannot run in real time)

Hard to optimize each component

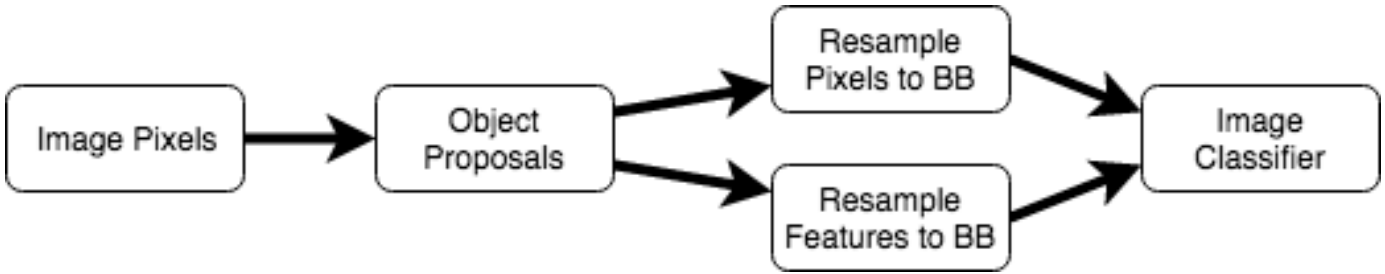

Yolo: You Only Look Once

Consider detection a regression problem

Use a single ConvNet

Runs once on entire image. Very Fast!

You only look once: Unified, real-time object detection.
Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi
June 2015

How it works

The following predictions are made for each cell in an S x S grid.

C conditional class probabilities **$Pr(Class_i | Obj)$**

B bounding boxes (4 parameters each)

B confidence scores **Pr(Obj)*IoU**

Output is **S x S x (5B+C)** tensor

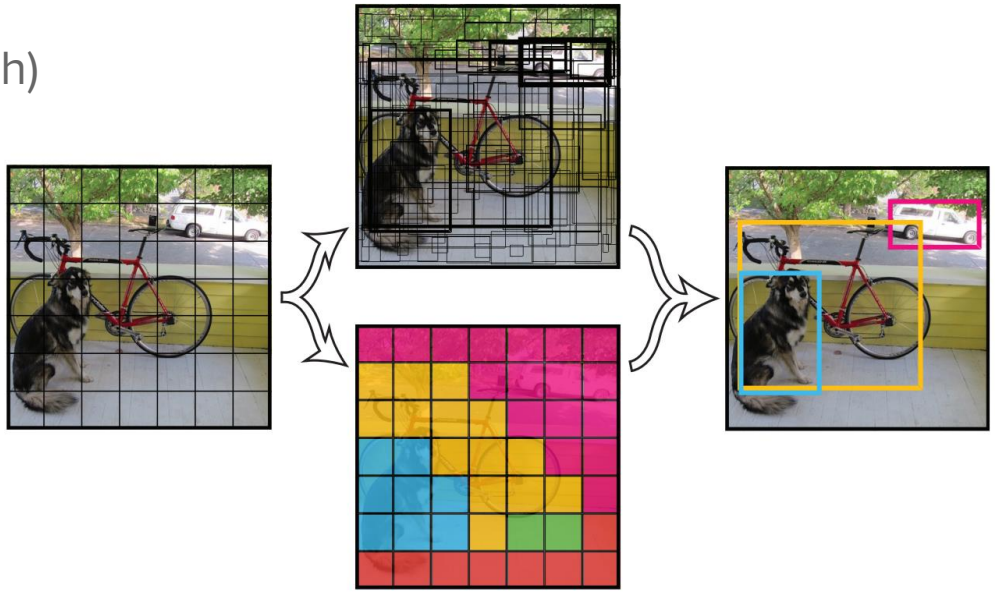

https://arxiv.org/pdf/1506.02640v5.pdf

Architecture

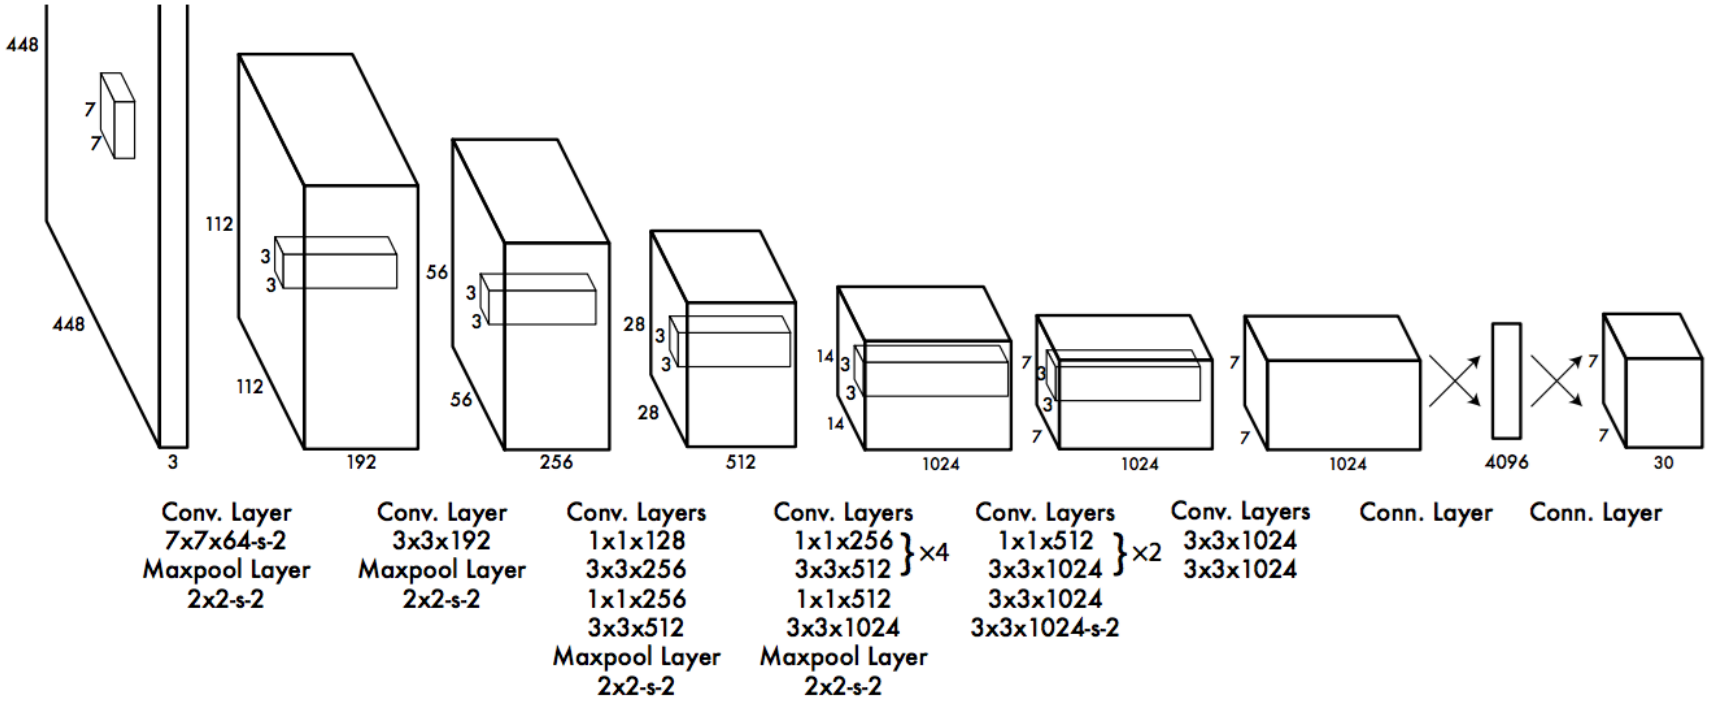

https://arxiv.org/pdf/1506.02640v5.pdf

Architecture

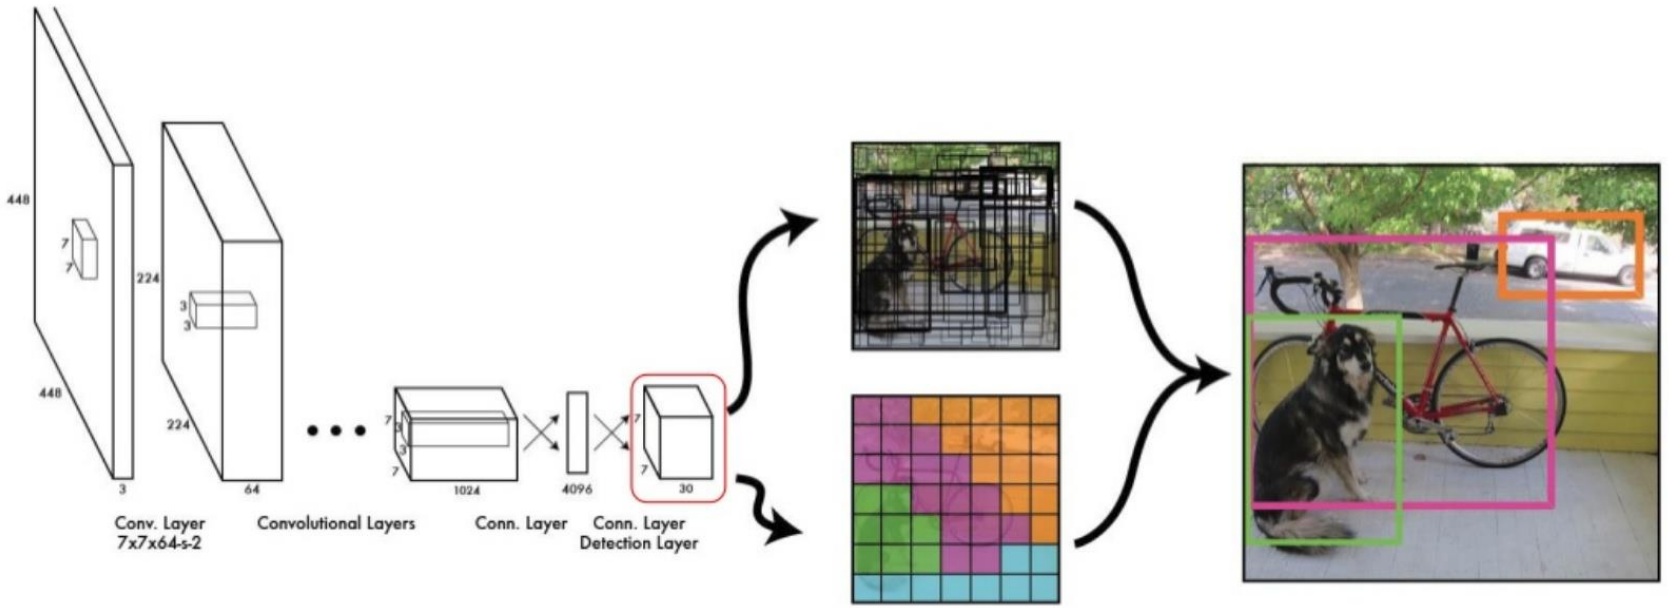

$$\Pr(\text{Class}_i|\text{Object}) * \Pr(\text{Object}) * \text{IOU}_{\text{pred}}^{\text{truth}} = \Pr(\text{Class}_i) * \text{IOU}_{\text{pred}}^{\text{truth}}$$

https://arxiv.org/pdf/1506.02640v5.pdf

Architecture

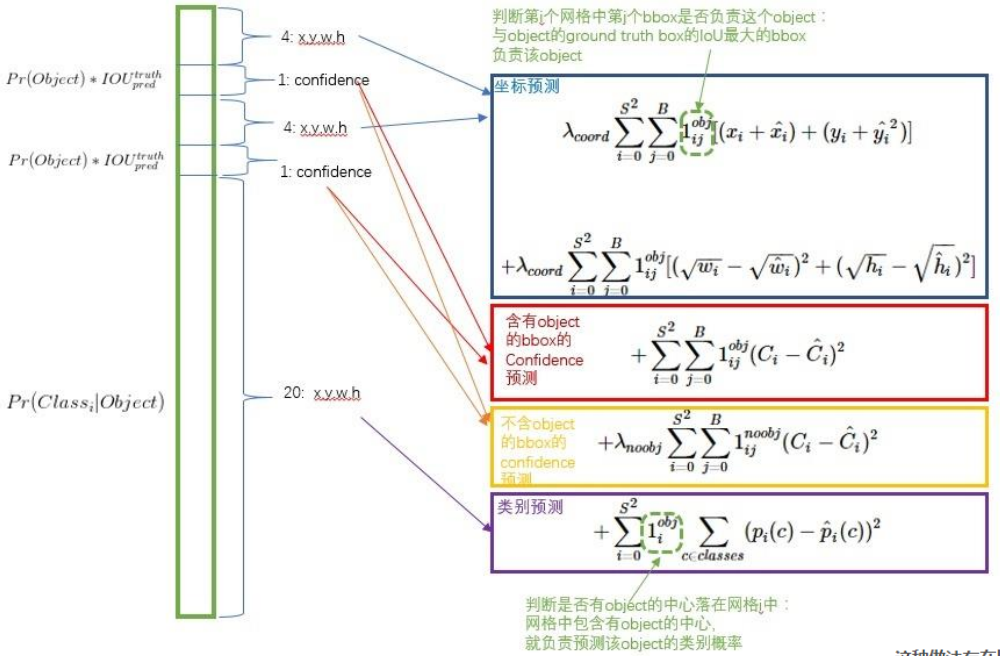

https://arxiv.org/pdf/1506.02640v5.pdf

这种做法存在以下几个问题：

第一，8维的localization error和20维的classification error同等重要显然是不合理的；

第二，如果一个网格中没有object（一幅图中这种网格很多），那么就会将这些网格中的box的confidence push到0，相比于较少的有object的网格，这种做法是overpowering的，这会导致网络不稳定甚至发散。

解决办法：

- 更重视8维的坐标预测，给这些损失前面赋予更大的loss weight, 记为λ_{coord}在pascal VOC训练中取5。
- 对没有object的box的confidence loss，赋予小的loss weight，记为λ_{noobj}在pascal VOC训练中取0.5。
- 有object的box的confidence loss和类别的loss的loss weight正常取1。

Performance (VOC 2007)

	Yolo	Faster R-CNN (VGG-16)
mAP	63.4	73.2
FPS	45	7

Trained on Pascal VOC 2007 + 2012 dataset

https://arxiv.org/pdf/1506.02640v5.pdf

Limitations of Yolo

Struggles with small objects

Struggles with unusual aspect ratios

Poor localization

SSD Single Shot Detector

Faster than Yolo, as accurate as Faster R-CNN

Predicts categories and box offsets

Uses small convolutional filters applied to feature maps

Makes predictions using feature maps of different scales

SSD: Single shot multibox detector

Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, Alexander C. Berg

Dec 2015

Comparison to Yolo

38*38*4+19*19*6+10*10*6+5*5*6+3*3*4+1*1*4=8732 default boxes

SSD

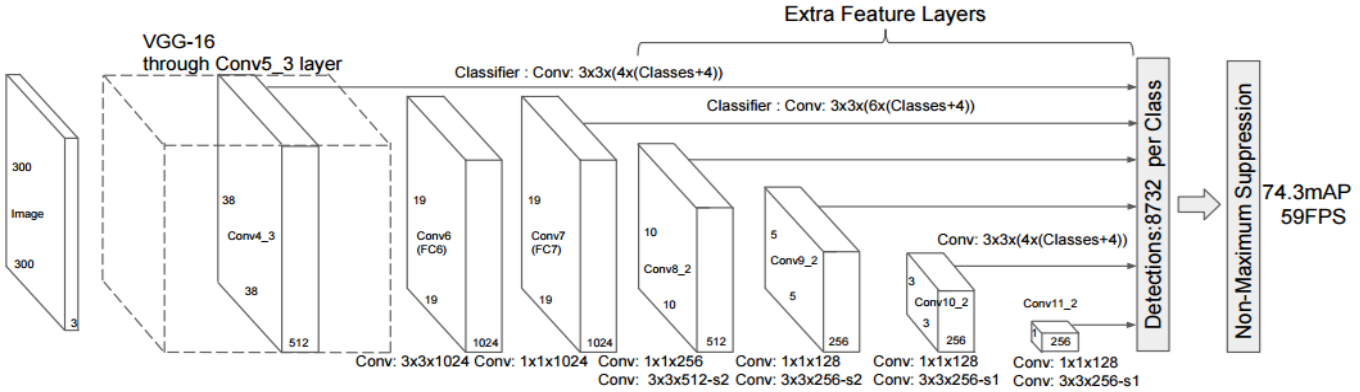

Yolo

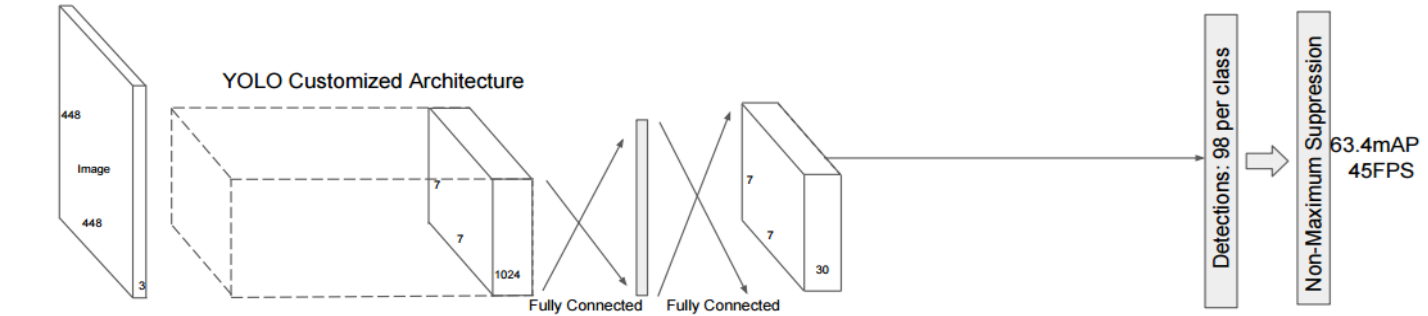

https://arxiv.org/pdf/1512.02325v5.pdf https://zhuanlan.zhihu.com/p/33544892

Default Boxes

Multiple aspect ratios per cell.

Similar to Faster R-CNN Anchor Boxes.

- Applied to many feature maps.

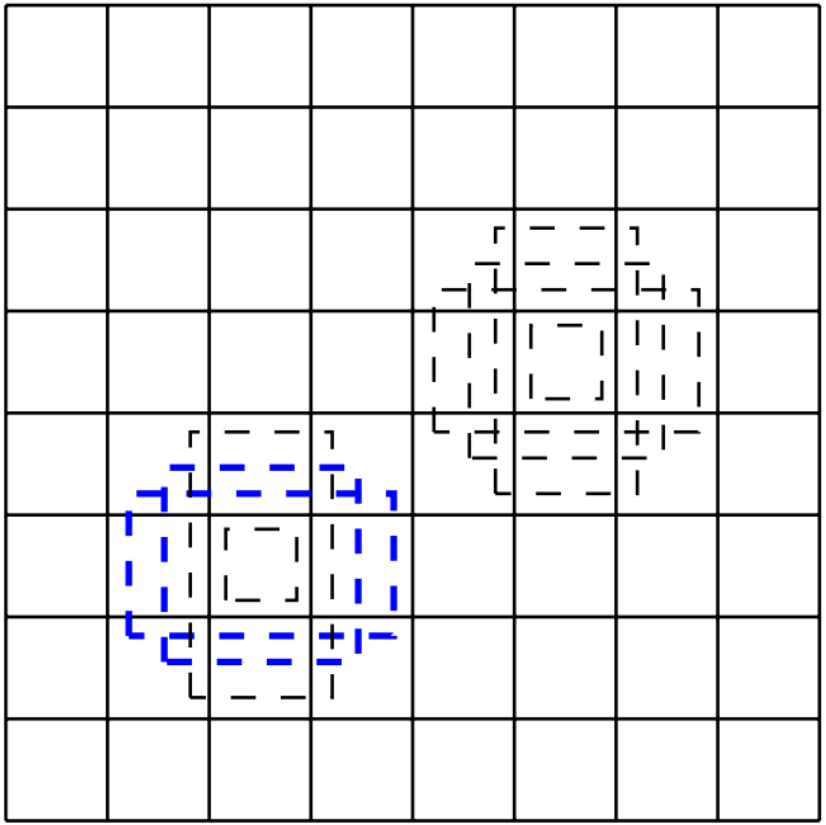

https://arxiv.org/pdf/1512.02325v5.pdf

SSD Detector

Detectors are convolutional filters.

Each detector outputs a single value.

Predict class probabilities.

Predict bounding box offsets.

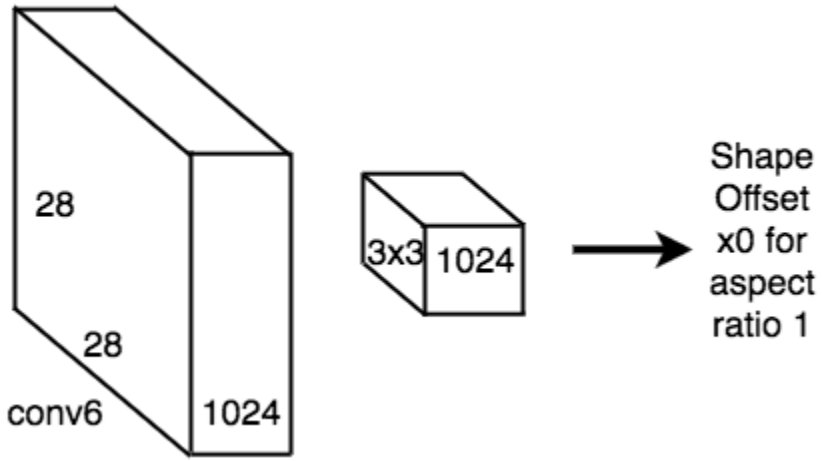

(classes + 4) detectors are needed for a detection.

(classes + 4) x (#default boxes) x m x n outputs for a mxn feature map.

Results (VOC 2007)

	SSD*	Yolo*	Faster R-CNN*
mAP	74.3	66.4	73.2
FPS	46	21	7

*VGG16

Trained on Pascal VOC 2007 + 2012 dataset

https://arxiv.org/pdf/1512.02325v5.pdf

Others

Part III

Feature Pyramid Networks

Uses ConvNet as Feature Pyramid

Includes low level feature maps to detect small objects

Top down pathway provides contextual information

Feature Pyramid Networks for Object Detection.

Tsung-Yi Lin, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, Serge Belongie.

Dec 2016

Comparison to prior methods

a. Accurate but slow.

b. Misses low level information. (Yolo)

c. Misses context in low level predictions. (SSD)

d. Accurate and fast.

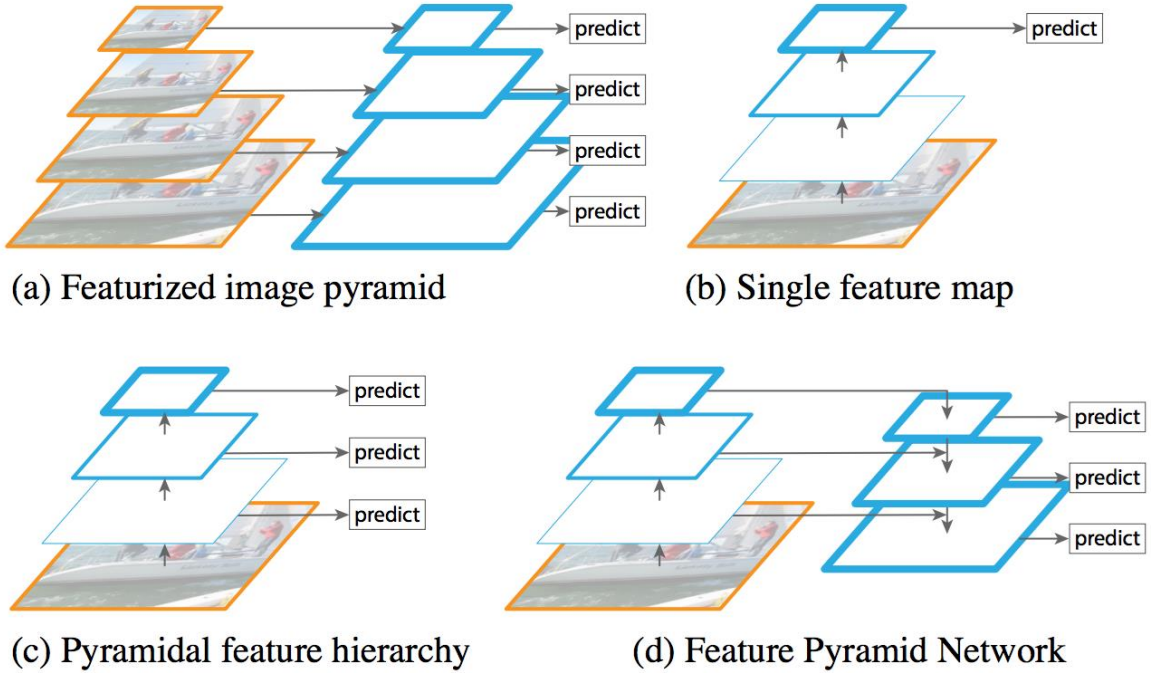

https://arxiv.org/pdf/1512.02325v5.pdf

Top down pathway

- High level (semantically strong) feature maps are upsampled.
- Lateral connections merge feature maps from bottom up pathway.

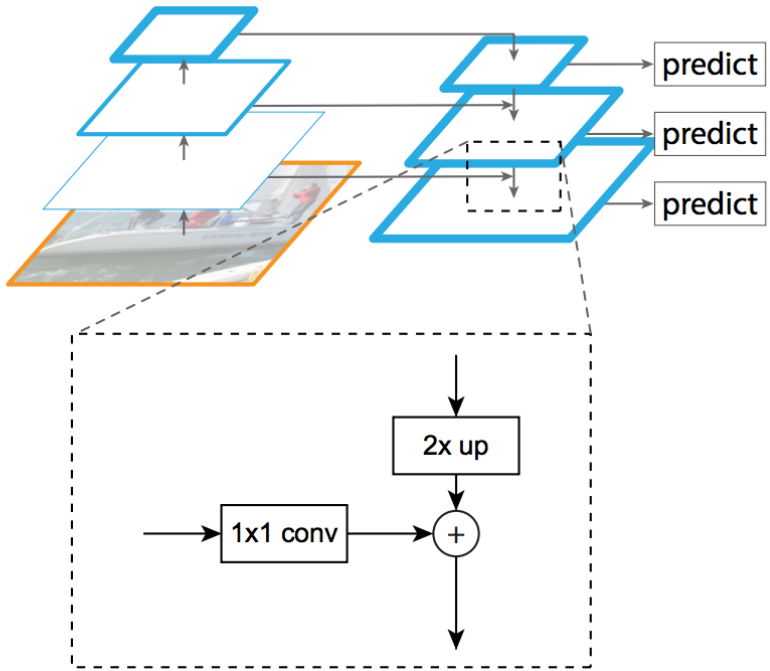

https://arxiv.org/pdf/1512.02325v5.pdf

Feature Pyramid Networks for RPN

- Replace single scale feature map with FPN.
- Single scale anchors at each level.

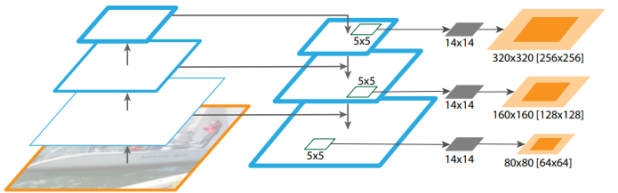

Figure 4. FPN for object segment proposals. The feature pyramid is constructed with identical structure as for object detection. We apply a small MLP on 5×5 windows to generate dense object segments with output dimension of 14×14. Shown in orange are the size of the image regions the mask corresponds to for each pyramid level (levels P_{3-5} are shown here). Both the corresponding image region size (light orange) and canonical object size (dark orange) are shown. Half octaves are handled by an MLP on 7x7 windows ($7 \approx 5\sqrt{2}$), not shown here. Details are in the appendix.

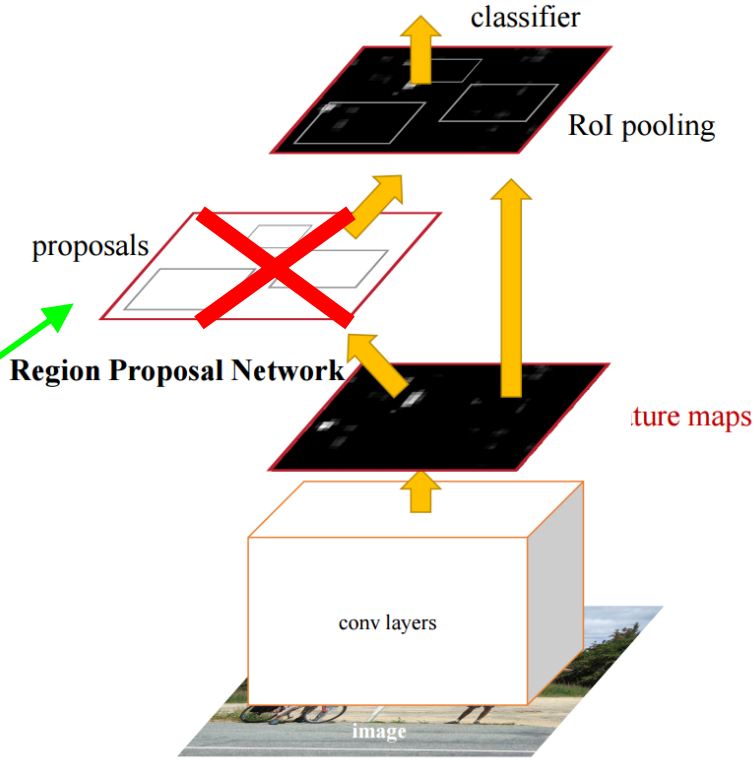

https://arxiv.org/abs/1612.03144

Results

- COCO

- Previous State of the art single-model

	Faster R-CNN on **FPN***	Faster R-CNN +++*	ION**
mAP	36.2	34.9	31.2
mAP (small images)	18.2	15.6	12.8

* ResNet-101
** VGG-16

https://arxiv.org/abs/1612.03144

ION : Inside-Outside Network

Use multi-scale Conv features for inside region

Use four direction RNN features for outside region

Inside-Outside Net: Detecting Objects in Context with Skip Pooling and Recurrent Neural Networks

Sean Bell, C.Lawrence Zitnick, Kavita Bala, Ross Girshick
Cornell University, Microsoft Research

CVPR 2016

ION: Inside-Outside Net

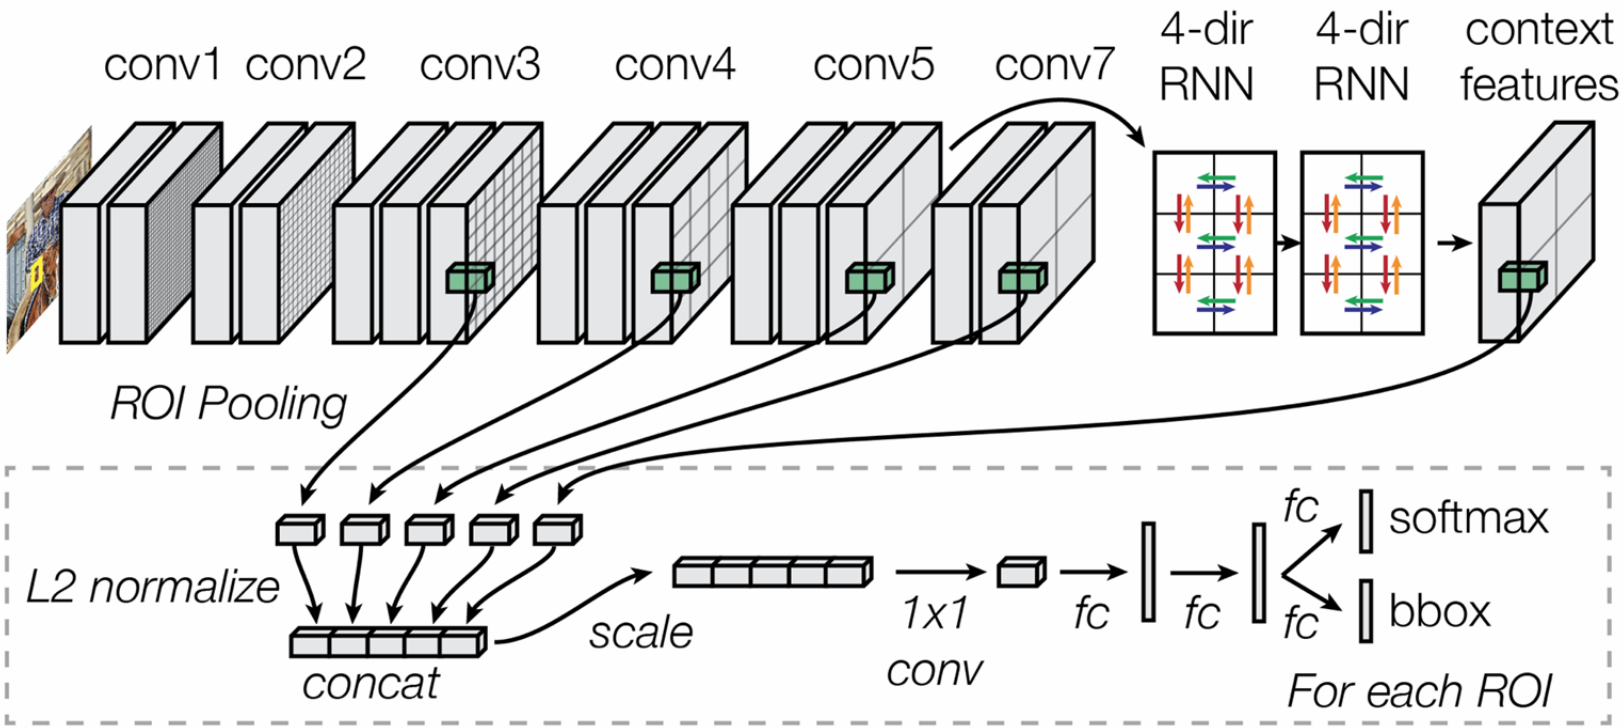

https://www.robots.ox.ac.uk/~vgg/rg/slides/ion-coco.pdf

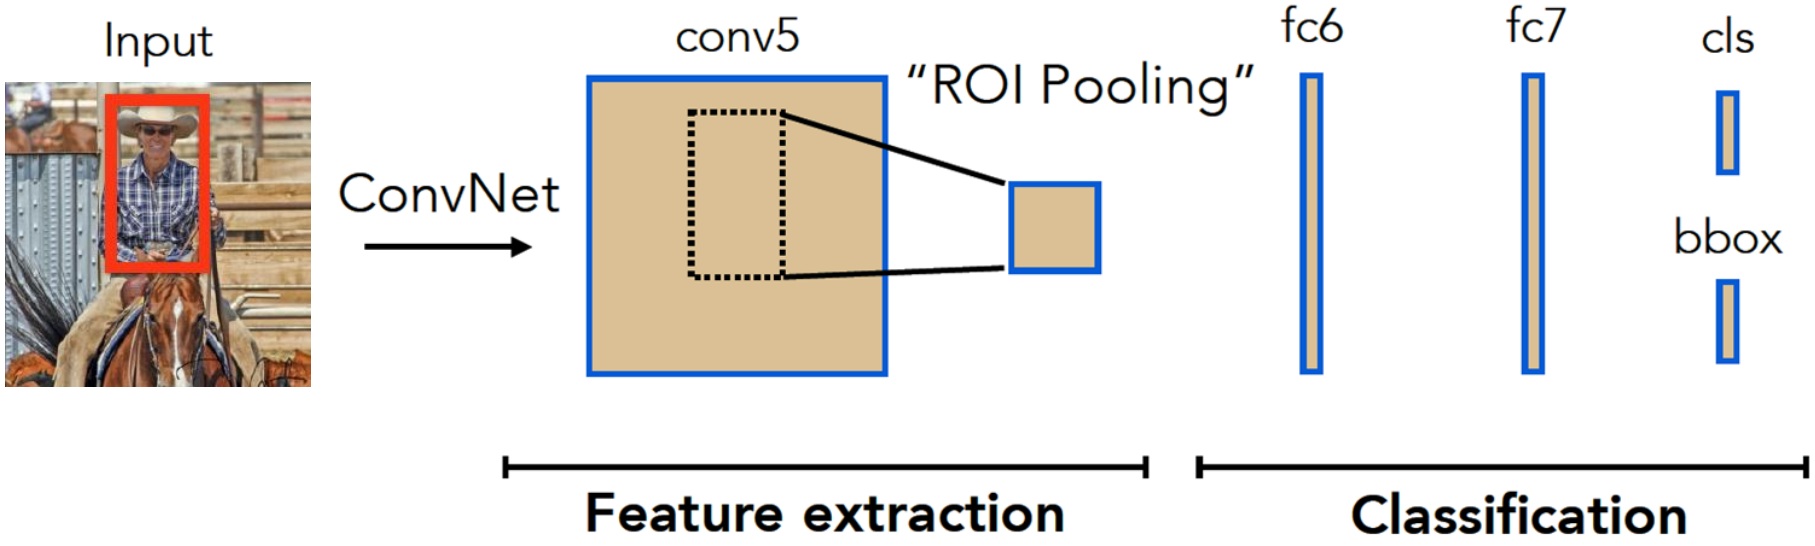

Can we improve on feature extraction?

- For small objects, the footprint on conv5 might only cover a 1x1 cell, which gets upsampled to 7x7
- Only local features (inside the ROI) are used for classification

COMBINING ACROSS LAYERS

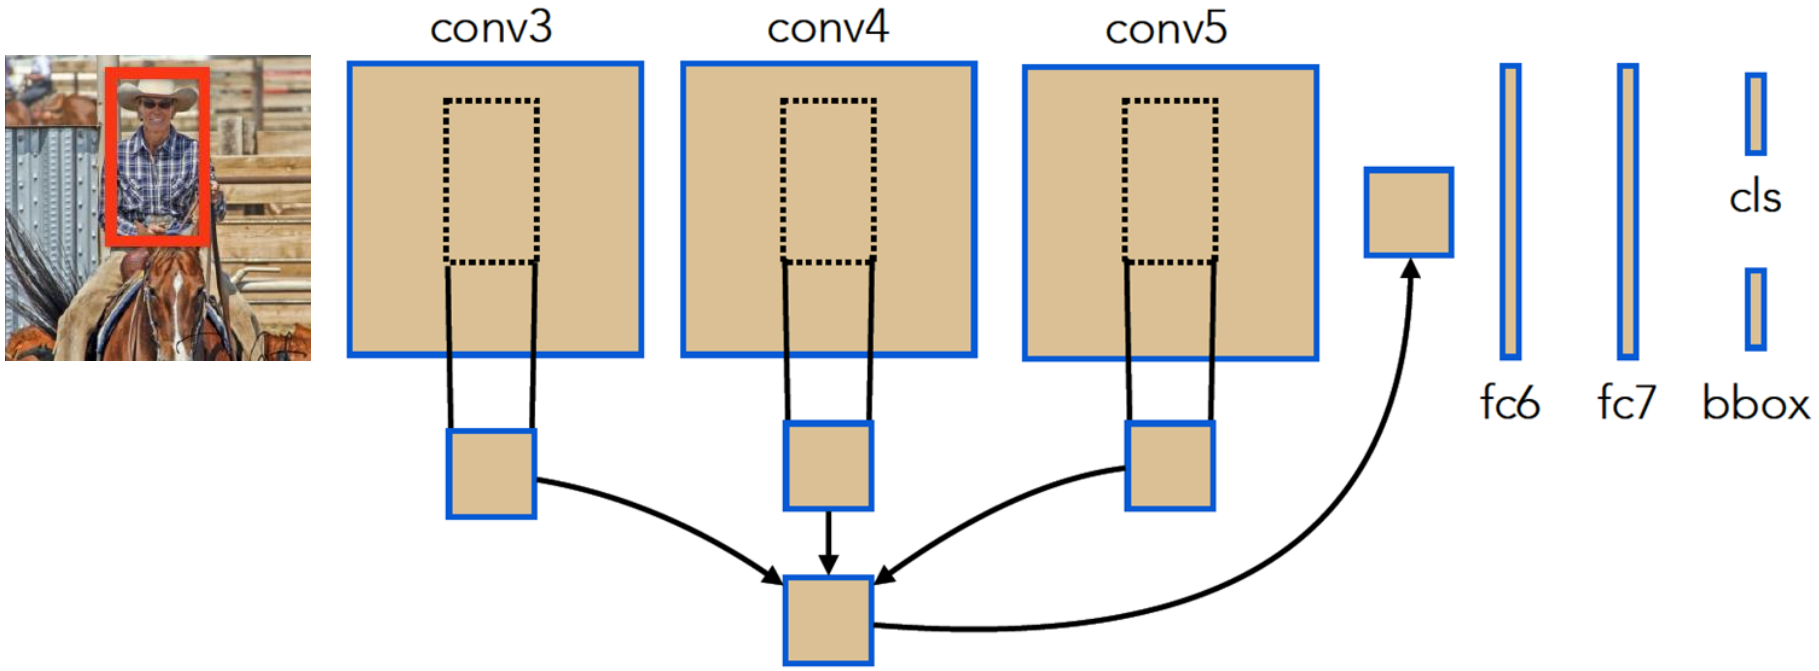

normalize, concatenate, **re-scale**

Uses RNNs to capture context info from outside bounding box.

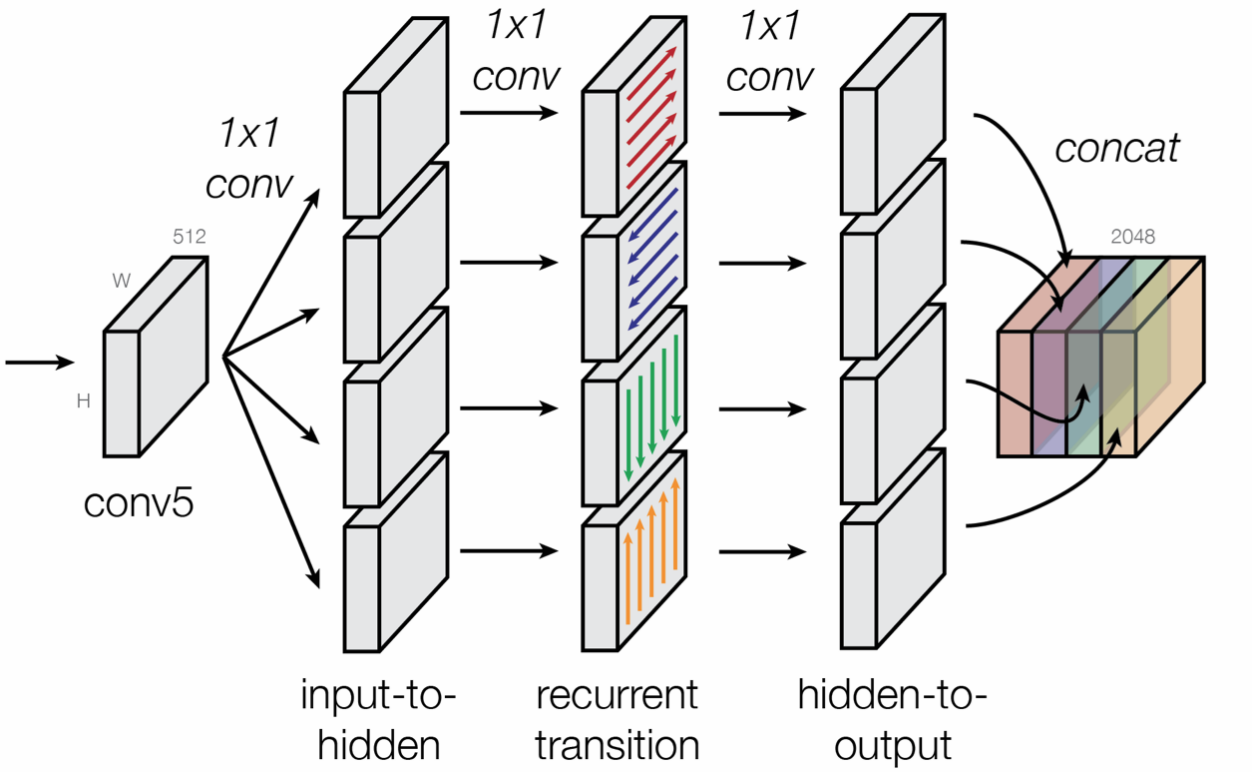

Stack 2 RNNs together

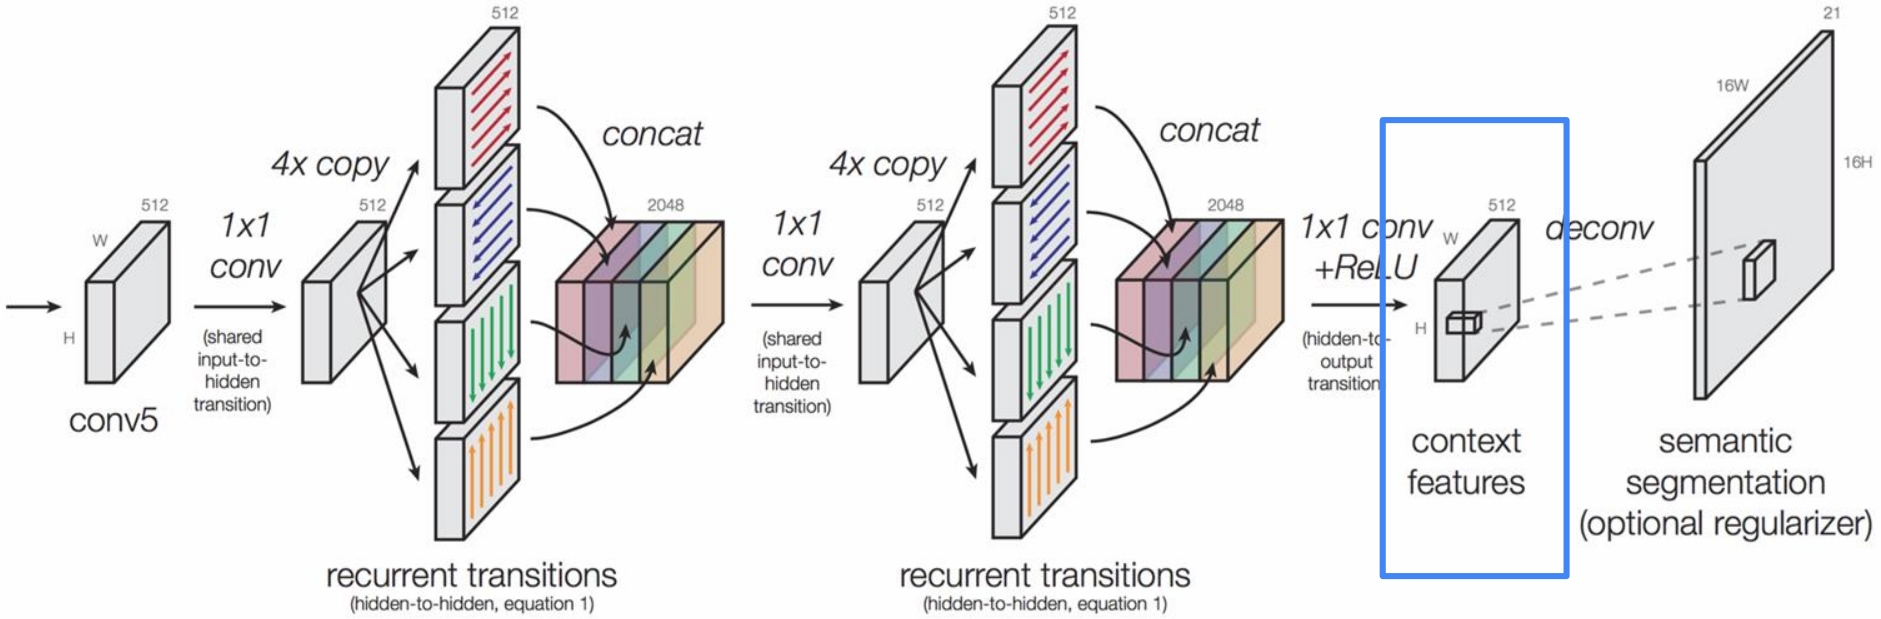

ION: Inside-Outside Net

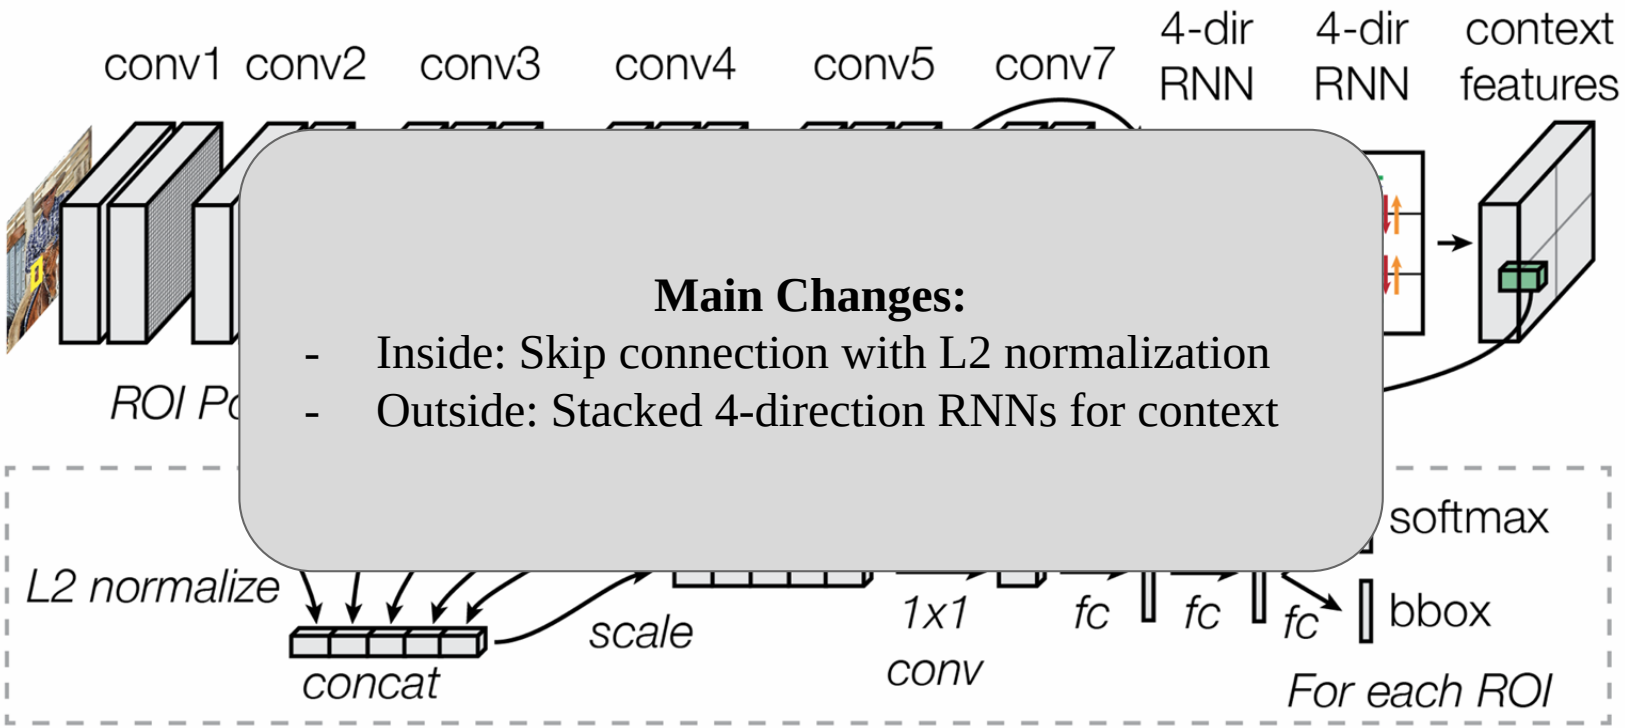

Main Changes:

- Inside: Skip connection with L2 normalization
- Outside: Stacked 4-direction RNNs for context

Results (VOC 2007)

METHOD	MAP
FAST R-CNN [GIRSHICK 2014]	70.0
FASTER R-CNN [GIRSHICK 2015]	73.2
CONV3+CONV4+CONV5	**75.6**
+ RNN + SEGMENTATION LOSS	**76.5**
+ SECOND BBOX REGRESSION + WEIGHTED VOTING	**78.5**
— DROPOUT	**79.2**

https://www.robots.ox.ac.uk/~vgg/rg/slides/ion-coco.pdf

Trained on Pascal VOC 2007 + 2012 dataset

Summary and Comparison

Part IV

Speed / Accuracy Trade-off

Unified tensorflow architecture

Compare speed, accuracy and memory usage

Speed/Accuracy trade-offs for modern convolutional object detectors
Jonathan Huang, Kevin Murphy et al.
Nov 2016

https://people.eecs.berkeley.edu/~rbg/papers/r-cnn-cvpr.pdf

Same Architectures

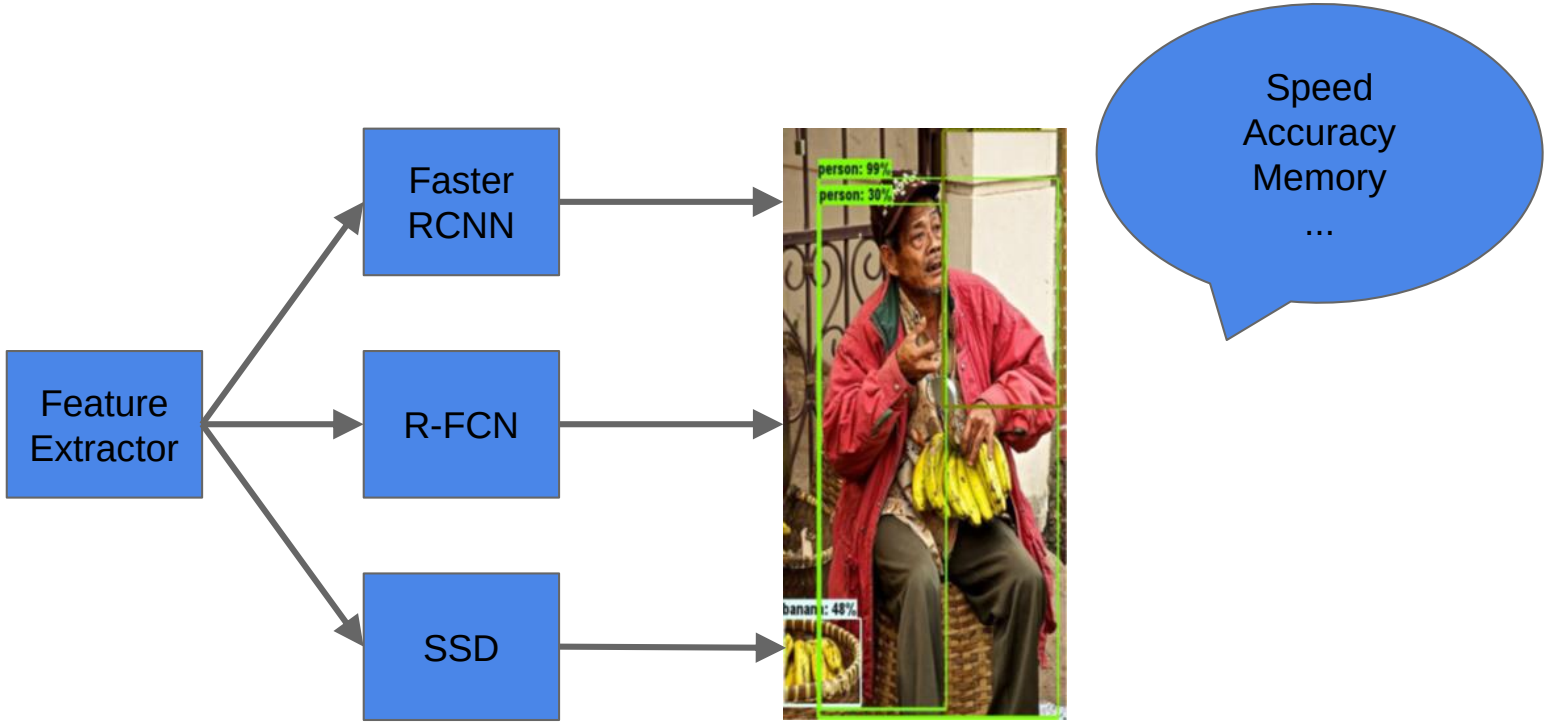

https://pan.baidu.com/s/1pKIKIIB

Recap: Faster RCNN

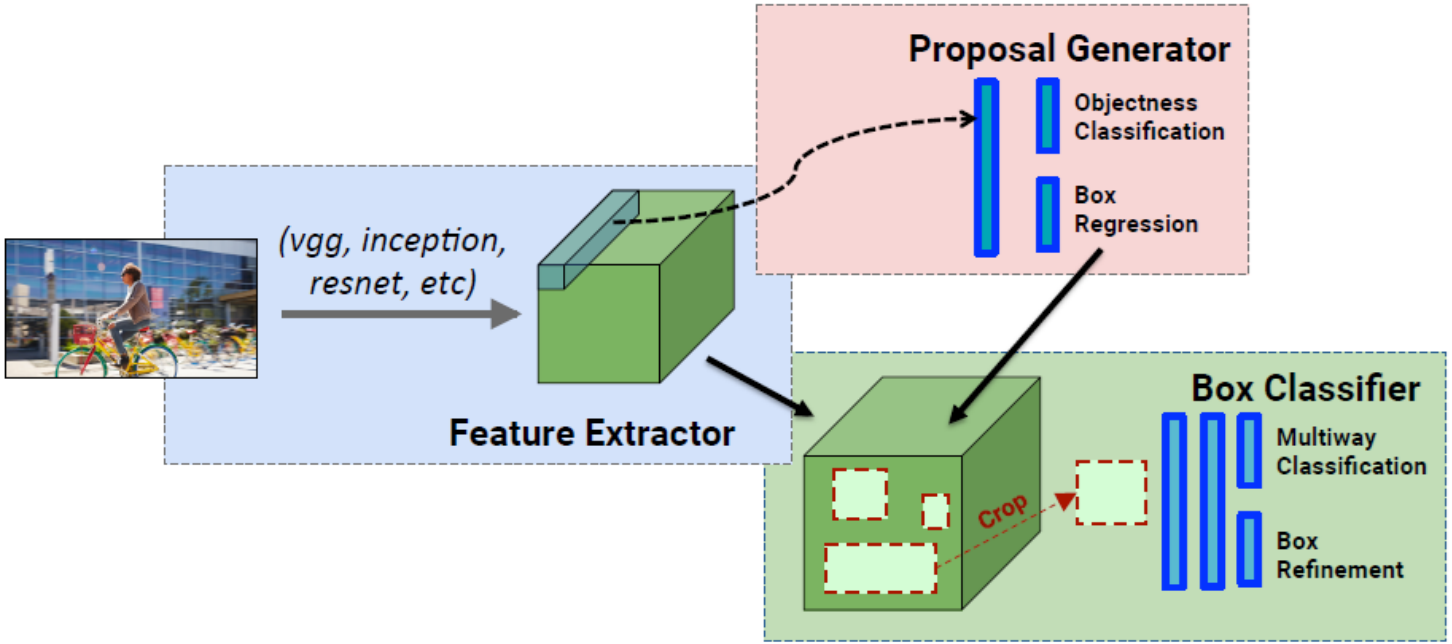

https://pan.baidu.com/s/1pKIKIIB

Recap: R-FCN

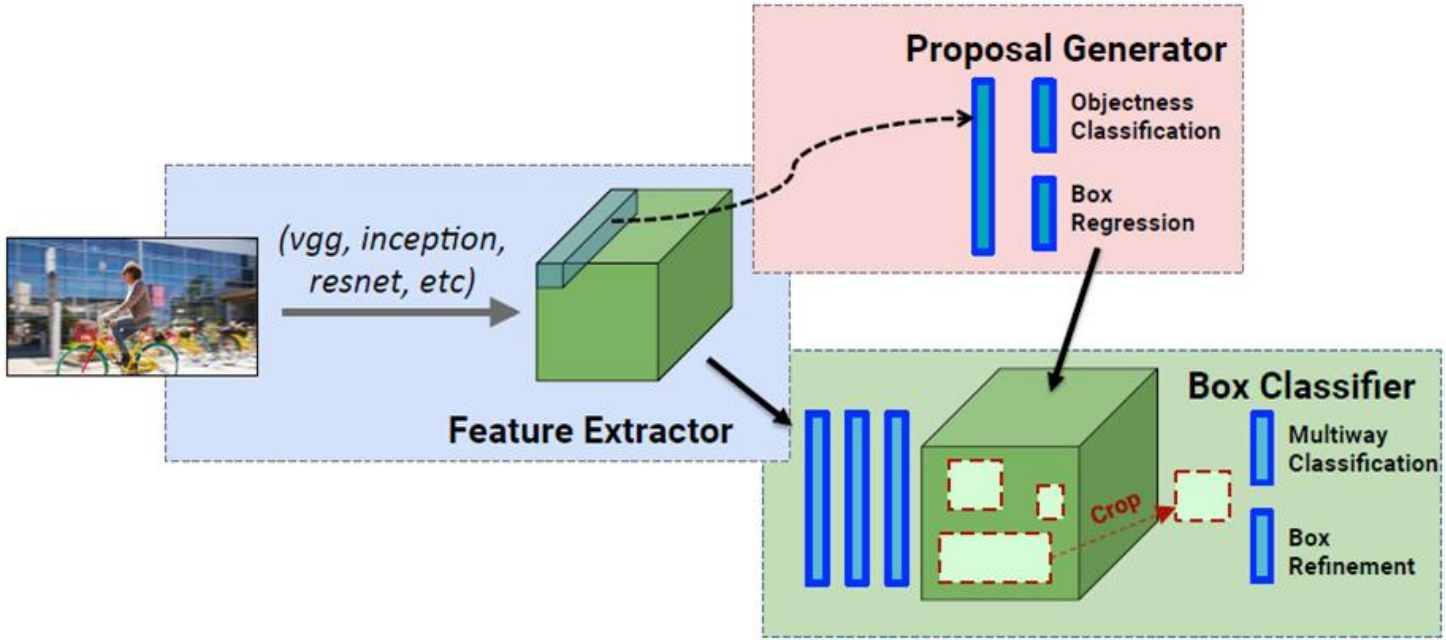

https://pan.baidu.com/s/1pKIKIIB

Recap: SSD

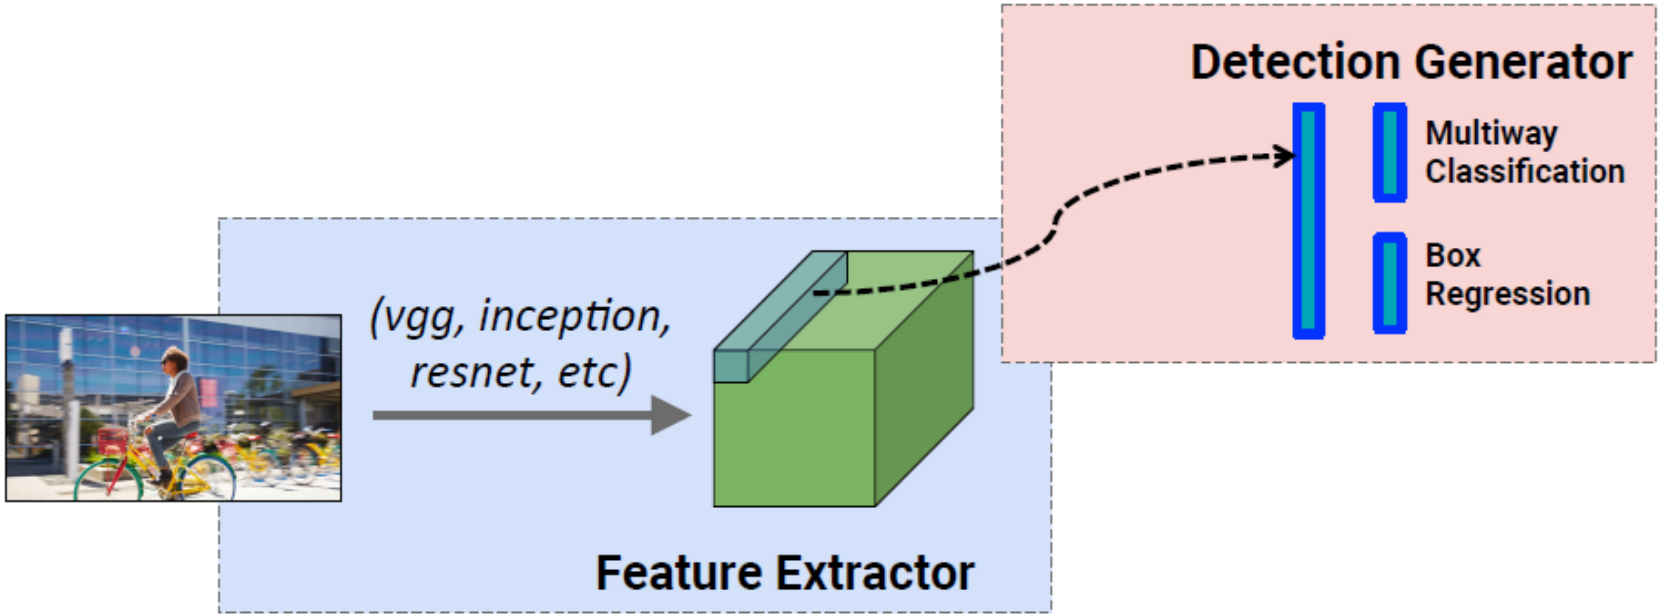

https://pan.baidu.com/s/1pKIKIIB

Accuracy VS Speed

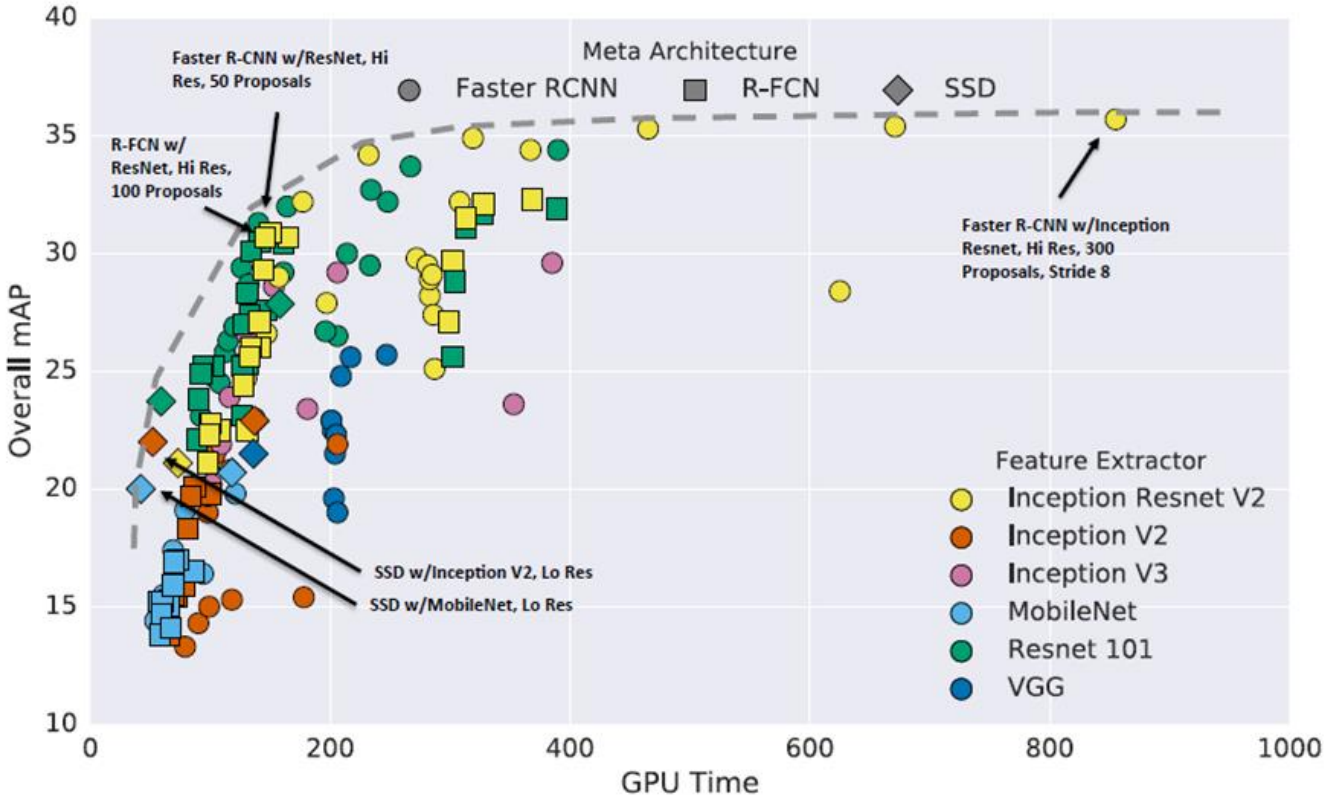

1. Different Feature Extractor

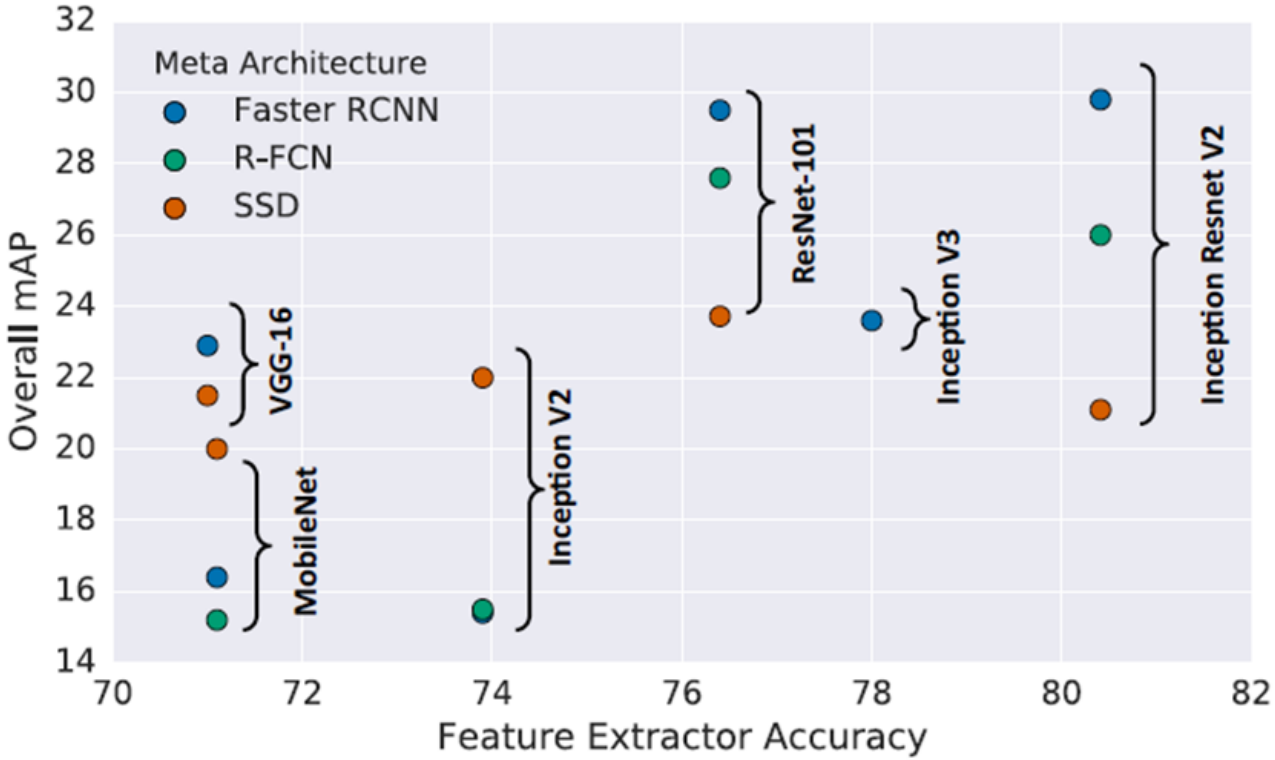

2. Detect Object Size

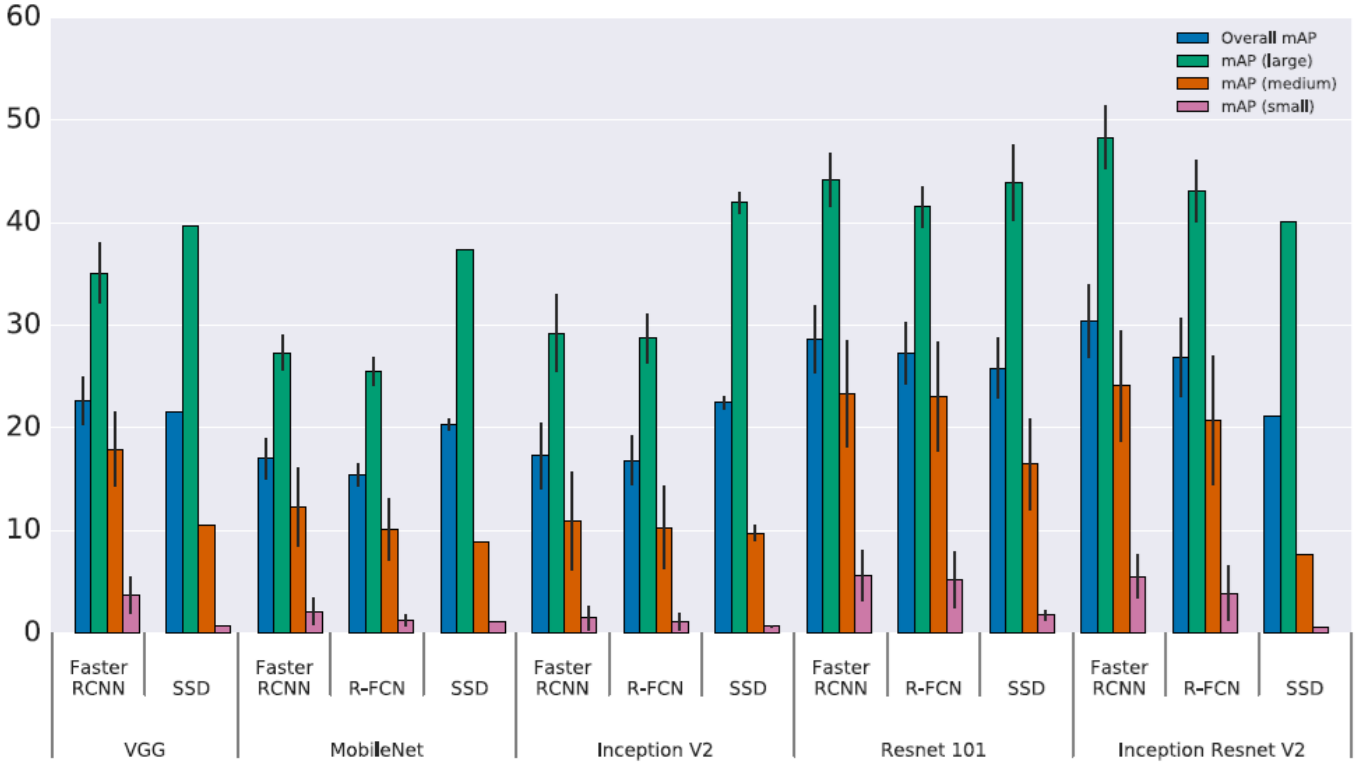

Figure 4: Accuracy stratified by object size, meta-architecture and feature extractor. We fix the image resolution to 300.

3. Image Resolution

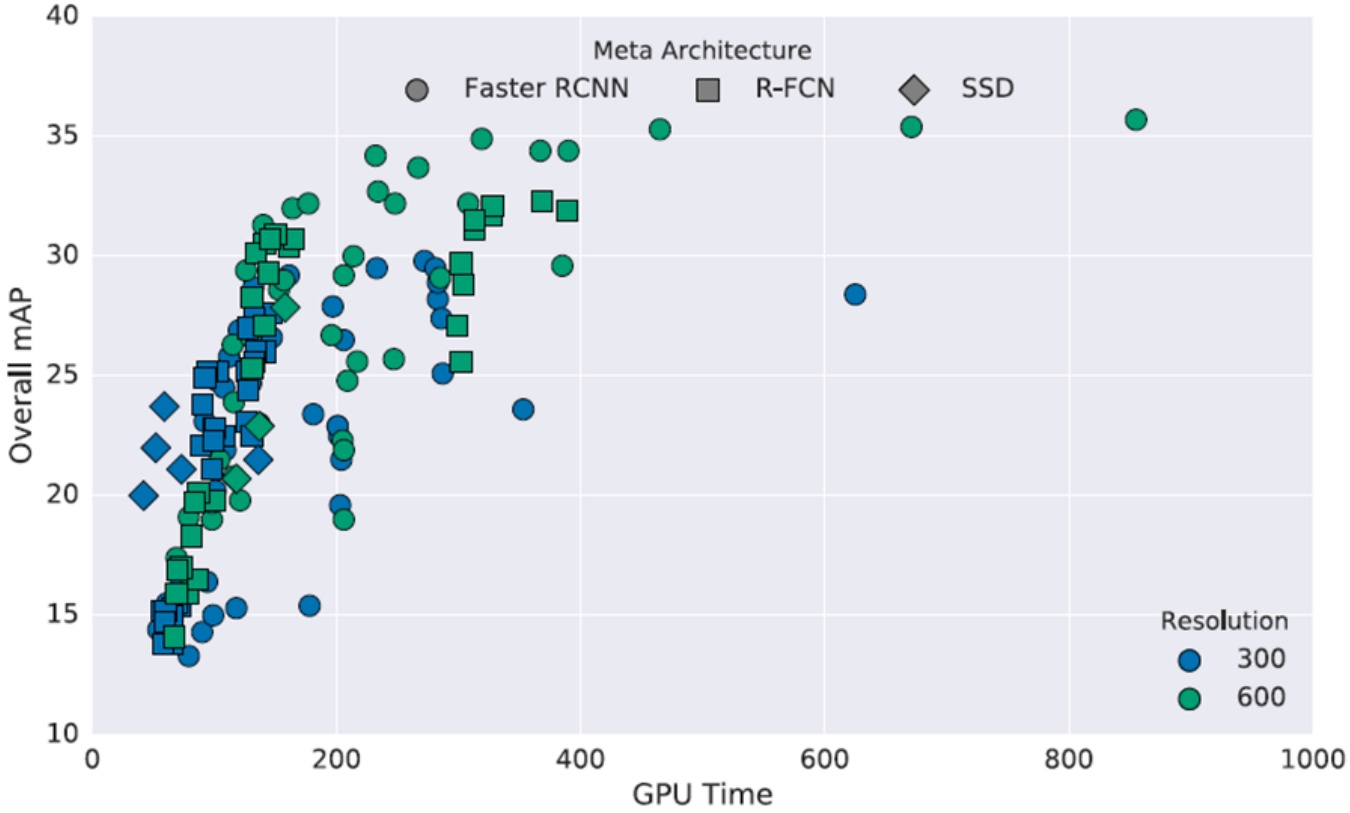

4. Region Proposal Number

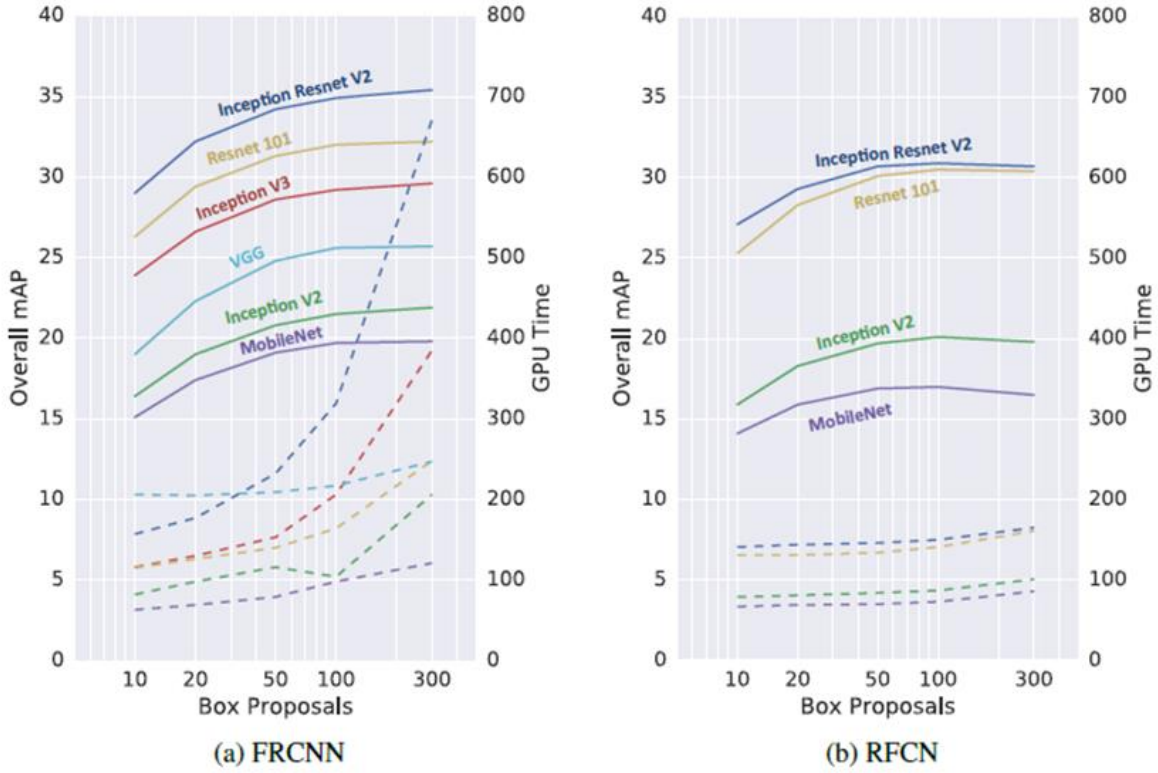

5. GPU Time

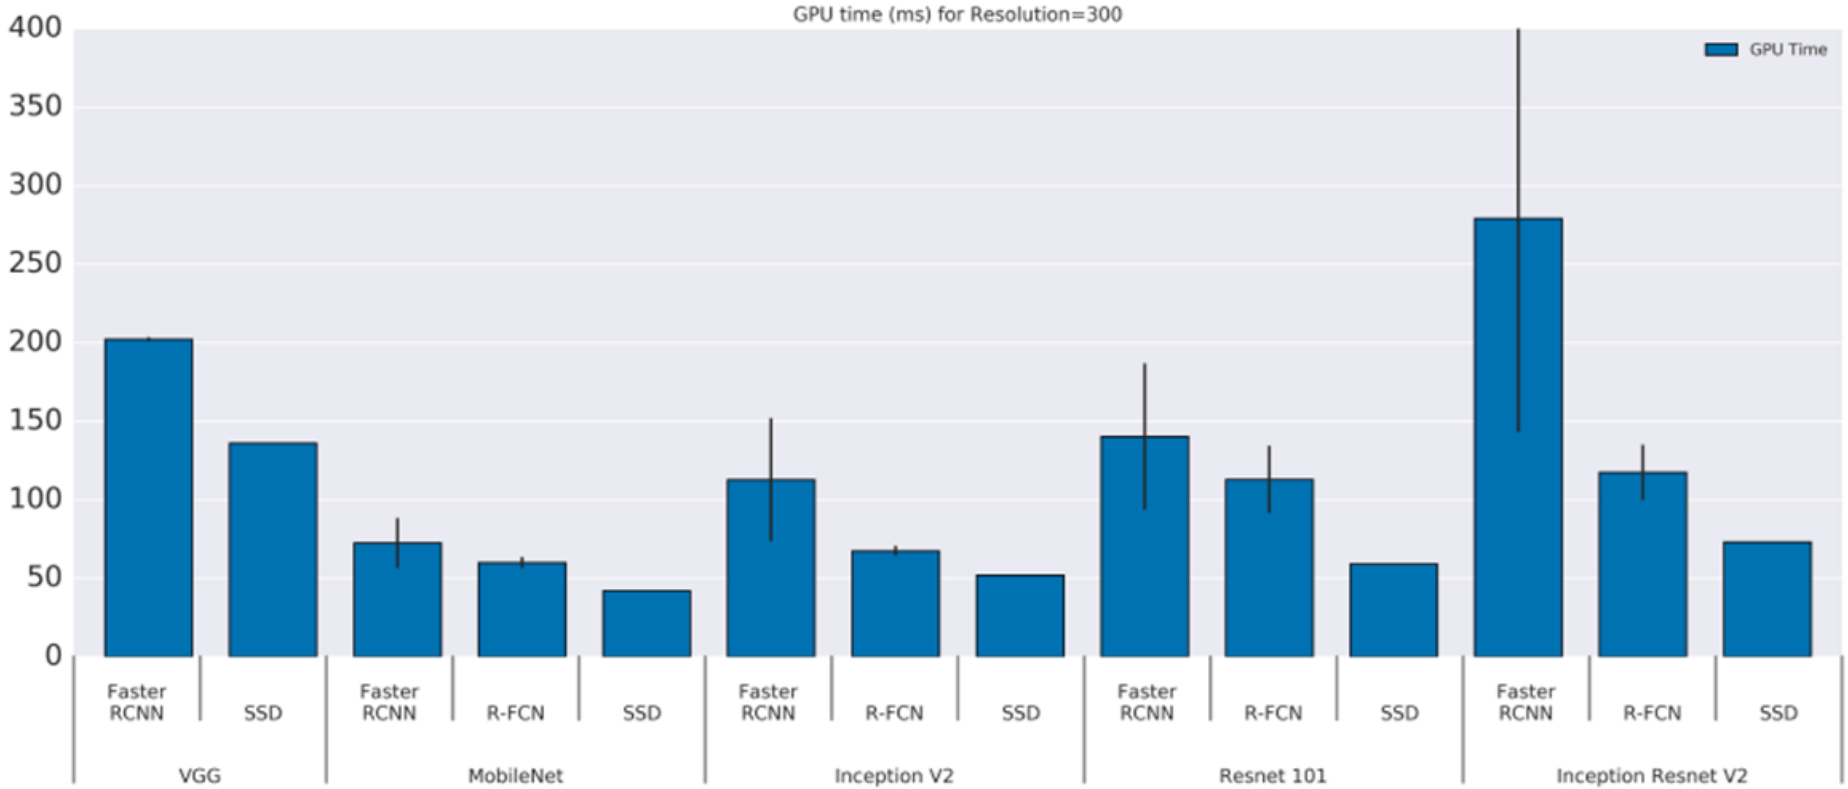
GPU time (ms) for Resolution=300
GPU Time
400
350
300
250
200
150
100
50
0
Faster RCNN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
VGG
MobileNet
Inception V2
Resnet 101
Inception Resnet V2

6. Memory

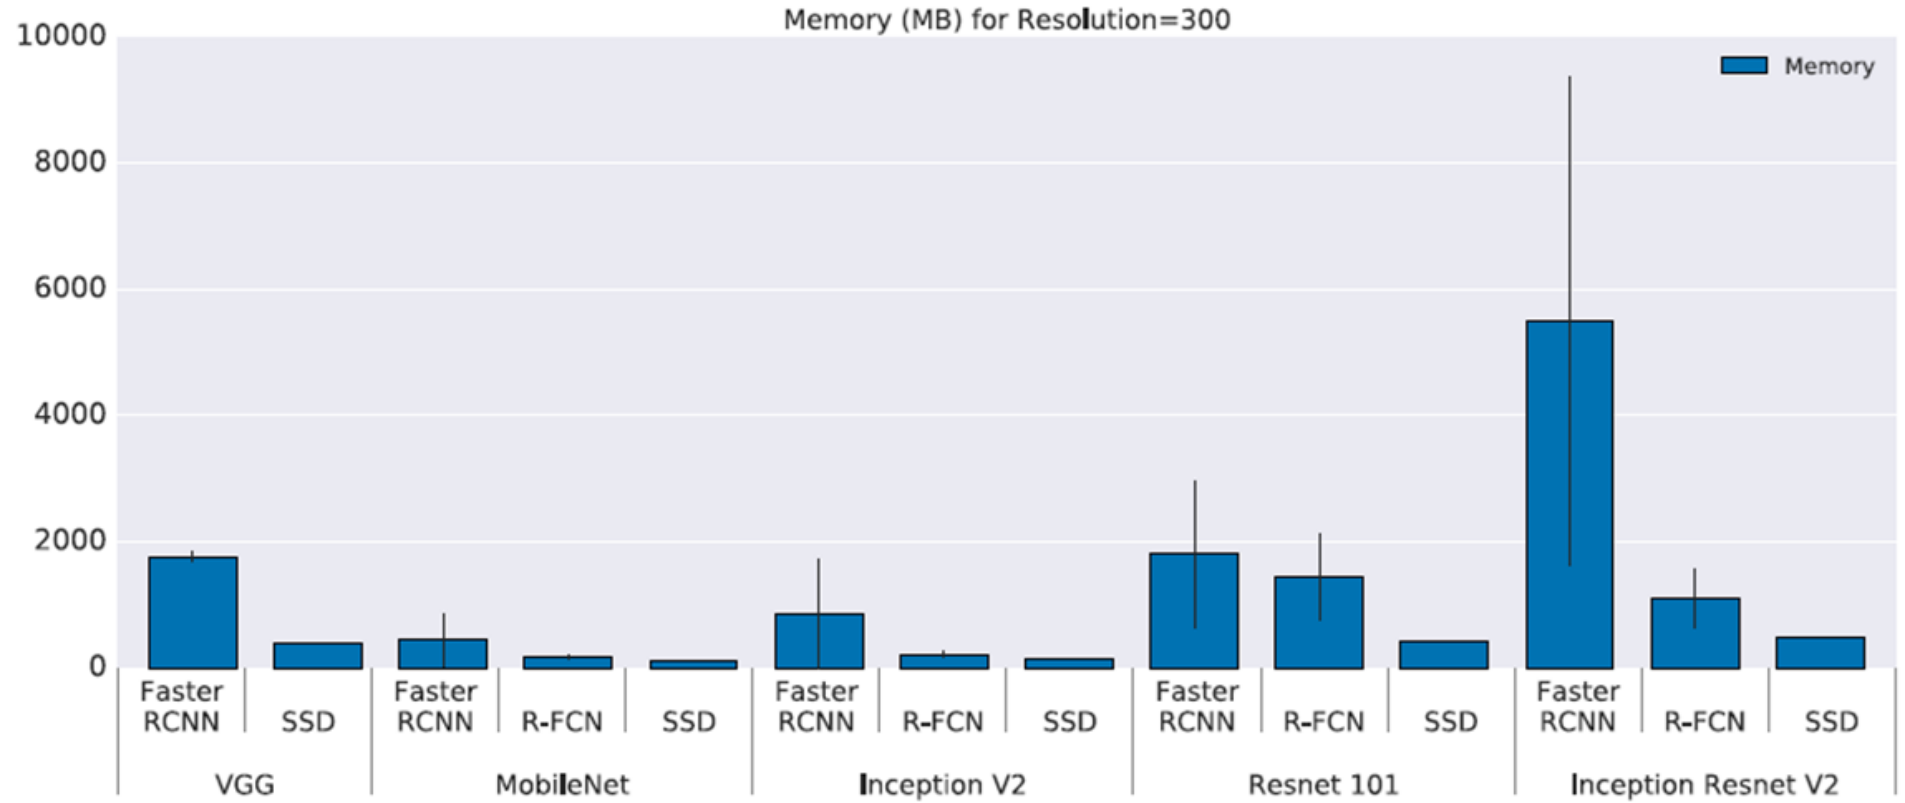
Memory (MB) for Resolution=300
Memory
10000
8000
6000
4000
2000
0
Faster RCNN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
VGG
MobileNet
Inception V2
Resnet 101
Inception Resnet V2

Summary

Speed First: SSD

Balance Speed and Accuracy: R-FCN

Accuracy First: Faster RCNN (Reduce proposals, it can speed up a lot with some accuracy loss)

Thanks!

Questions?

References

Girshick, Ross and Donahue, Jeff and Darrell, Trevor and Malik, Jitendra, Rich feature hierarchies for accurate object detection and semantic segmentation, CVPR 2014

He, Kaiming and Zhang, Xiangyu and Ren, Shaoqing and Sun, Jian, Spatial Pyramid Pooling in Deep Convolutional Networks for Visual Recognition, ECCV 2014

Girshick, Ross, Fast R-CNN, ICCV 2015

Ren, Shaoqing and He, Kaiming and Girshick, Ross and Sun, Jian, Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks, CVPR 2015

Jifeng Dai, Yi Li, Kaiming He, Jian Sun R-FCN: Object Detection via Region-based Fully Convolutional Networks, NIPS 2016

Erhan, Dumitru and Szegedy, Christian and Toshev, Alexander and Anguelov, Dragomir, Scalable Object Detection using Deep Neural Networks, CVPR 2014

Bell, Sean and Lawrence Zitnick, C and Bala, Kavita and Girshick, Ross, Inside-Outside Net: Detecting Objects in Context with Skip Pooling and Recurrent Neural Networks, CVPR 2016

Redmon, Joseph and Divvala, Santosh and Girshick, Ross and Farhadi, Ali, You Only Look Once: Unified, Real-Time Object Detection, CVPR 2016

Liu, Wei and Anguelov, Dragomir and Erhan, Dumitru and Szegedy, Christian and Reed, Scott and Fu, Cheng-Yang and Berg, Alexander C, SSD: Single Shot MultiBox Detector, ECCV 2016

Tsung-Yi Lin, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, Serge Belongie, Feature Pyramid Networks for Object Detection, arXiv 2016

Huang, Jonathan and Rathod, Vivek and Sun, Chen and Zhu, Menglong and Korattikara, Anoop and Fathi, Alireza and Fischer, Ian and Wojna, Zbigniew and Song, Yang and